

# DSC 40B

Theoretical Foundations II

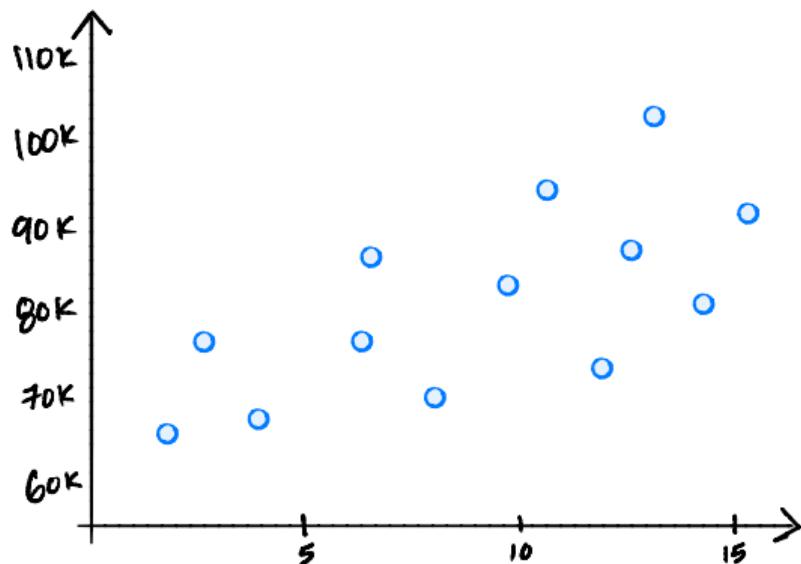
Lecture 1 | Part 1

**What is DSC 40B?**

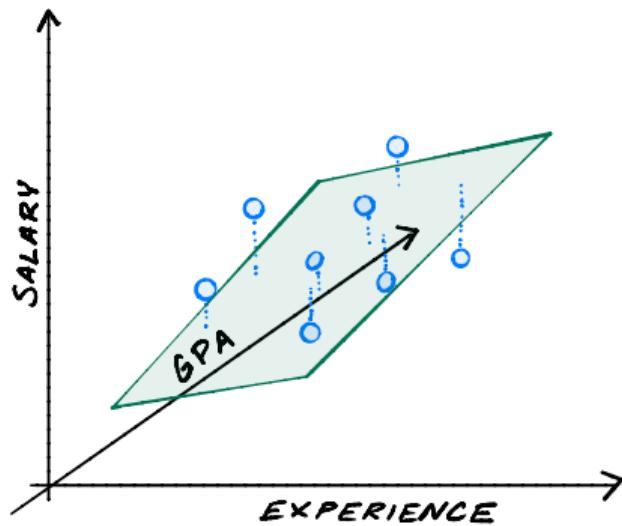
## Recall DSC 40A...

- ▶ How do we **formalize** learning from data?
- ▶ How do we turn it into something a **computer** can do?

# Example: Predicting Salary



# Example: Predicting Salary



# The End

$$(X^T X)^{-1} \vec{w} = X^T \vec{b}$$

# Wait...

- ▶ We actually need to **compute** the answer...
- ▶ We need an **algorithm**.

# An Algorithm?

```
>>> import numpy as np  
>>> w = np.linalg.solve(X.T @ X, X.T @ b)
```

- ▶ Will it work for 1,000,000 data points?
- ▶ What about for 1,000,000 features?

# Example: Minimize Error

- ▶ **Goal:** summarize a collection of numbers,  $x_1, \dots, x_n$ :
- ▶ **Idea:** find number  $M$  minimizing the total absolute error:

$$\sum_{i=1}^n |M - x_i|$$

## Example: Minimize Error

- ▶ **Solution:** The **median** of  $x_1, \dots, x_n$ .
- ▶ But how do we actually **compute** the median?

# DSC 40B

## Theoretical Foundations II

Lecture 1 | Part 2

Example: Clustering

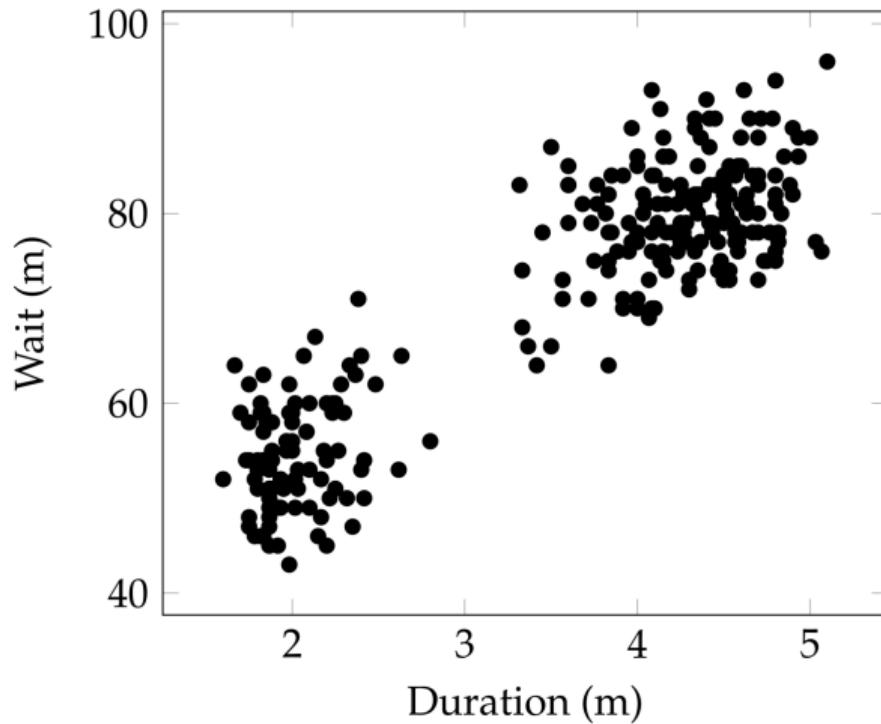
# Clustering

- ▶ Given a pile of data, discover similar groups.
- ▶ Examples:
  - ▶ Find political groups within social network data.
  - ▶ Given data on COVID-19 symptoms, discover groups that are affected differently.
  - ▶ Find the similar regions of an image (**segmentation**).
- ▶ Most useful when data is high dimensional...

# Example: Old Faithful



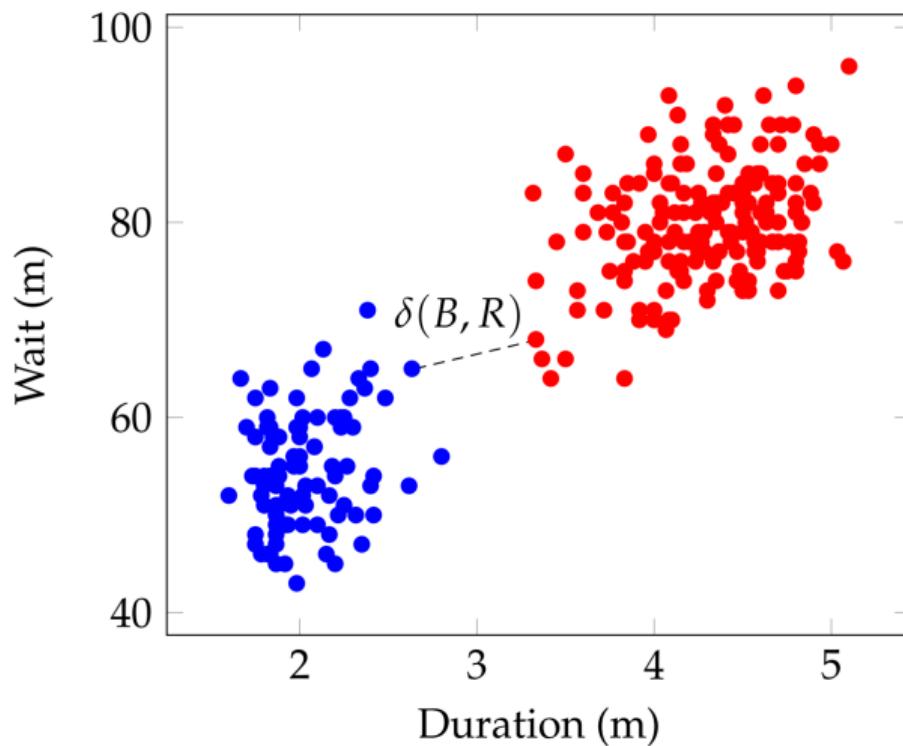
# Example: Old Faithful



# Clustering

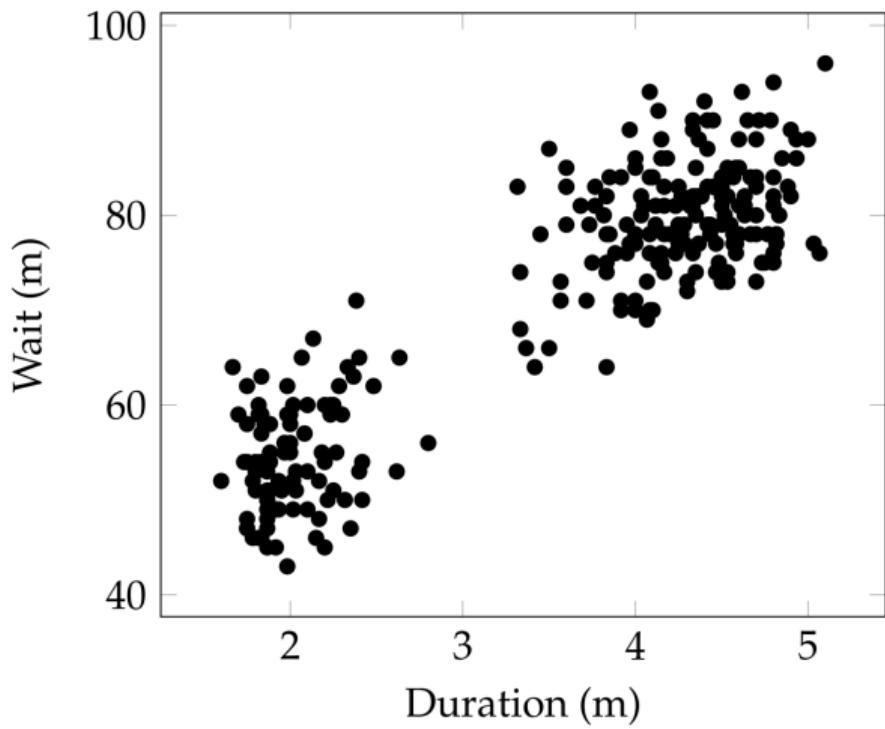
- ▶ Goal: for computer to identify the two groups in the data.

# Example: Old Faithful

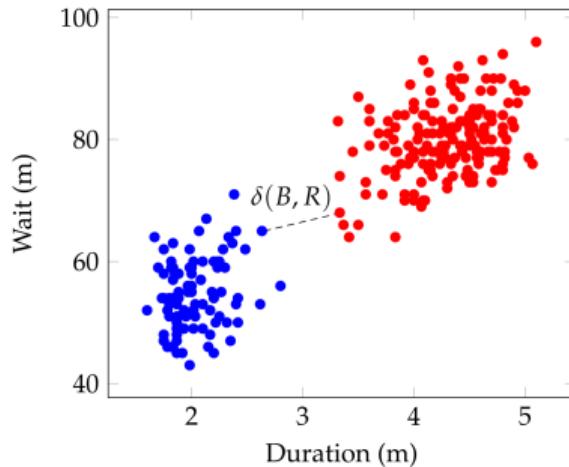


# Clustering

- ▶ How do we turn this into something a **computer** can do?
- ▶ DSC 40A says: “Turn it into an optimization problem”.
- ▶ **Idea:** develop a way of quantifying the “goodness” of a clustering; find the **best**.



# Quantifying Separation



Define the “separation”  $\delta(B, R)$  to be the smallest distance between a blue point and red point.

# The Problem

- ▶ **Given:**  $n$  points  $\vec{x}^{(1)}, \dots, \vec{x}^{(n)}$ .
- ▶ **Find:** an assignment of points to clusters  $\textcolor{red}{R}$  and  $\textcolor{blue}{B}$  so as to maximize  $\delta(\textcolor{blue}{B}, \textcolor{red}{R})$ .

# The End

# The “Brute Force” Algorithm

- ▶ There are finitely-many possible clusterings.
- ▶ **Algorithm:** Try each possible clustering, return that with largest separation,  $\delta(B, R)$ .
- ▶ This is called a **brute force** algorithm.

```
best_separation = float('inf') # Python for "infinity"
best_clustering = None

for clustering in all_clusterings(data):
    sep = calculate_separation(clustering)
    if sep < best_separation:
        best_separation = sep
        best_clustering = clustering

print(best_clustering)
```

# The End

# Wait...

- ▶ How long will this take to run if there are  $n$  points?
- ▶ How many clusterings of  $n$  things are there?

# Combinatorics

- ▶ How many ways are there to assign **R** or **B** to  $n$  objects?
- ▶ Two choices <sup>1</sup> for each object:  $2 \times 2 \times \dots \times 2 = 2^n$ .

---

<sup>1</sup>Small nitpick: actual color doesn't matter,  $2^{n-1}$ .

# Time

- ▶ Suppose it takes at least 1 nanosecond to check a single clustering.
  - ▶ One billionth of a second.
- ▶ If there are  $n$  points, it will take at least  $2^n$  nanoseconds to check all clusterings.

# Time Needed

| $n$ | Time          |
|-----|---------------|
| 1   | 1 nanosecond  |
| 10  | 1 microsecond |
| 20  | 1 millisecond |
| 30  | 1 second      |
| 40  | 18 minutes    |
| 50  | 13 days       |
| 60  | 36 years      |
| 70  | 37,000 years  |

# Example: Old Faithful

- ▶ The Old Faithful data set has 270 points.
- ▶ Brute force algorithm will finish in  $6 \times 10^{64}$  years.



# Algorithm Design

- ▶ Often, most obvious algorithm is **unusably slow**.
- ▶ Does this mean our problem is too hard?
- ▶ We'll see an efficient solution by the end of the quarter.

# DSC 40B

- ▶ Assess the efficiency of algorithms.
- ▶ Understand why and how common algorithms work.
- ▶ Develop faster algorithms using design strategies and data structures.

# **DSC 40B**

## *Theoretical Foundations II*

Lecture 1 | Part 3

**Measuring Efficiency by Timing**

# Efficiency

- ▶ Speed matters, *especially* with large data sets.
- ▶ An algorithm is only useful if it runs **fast enough**.
  - ▶ That depends on the size of your data set.
- ▶ How do we measure the efficiency of code?
- ▶ How do we know if a method will be fast enough?

# Scenario

- ▶ You're building a least squares regression model to predict a patient's blood oxygen level.
- ▶ You've trained it on 1,000 people.
- ▶ You have a full data set of 100,000 people.
- ▶ How long will it take? How does it **scale**?

# Example: Scaling

- ▶ Your code takes 5 seconds on 1,000 points.
- ▶ How long will it take on 100,000 data points?
- ▶  $5 \text{ seconds} \times 100 = 500 \text{ seconds?}$
- ▶ More? Less?

# Coming Up

- ▶ We'll answer this in coming lectures.
- ▶ Today: start with simpler algorithms for the mean, median.

# Approach #1: Timing

- ▶ How do we measure the efficiency of code?
- ▶ Simple: time it!
- ▶ Useful Jupyter tools: `time` and `timeit`

# Disadvantages of Timing

1. Time depends on the computer.
2. Depends on the particular input, too.
3. One timing doesn't tell us how algorithm **scales**.

# DSC 40B

## Theoretical Foundations II

Lecture 1 | Part 4

### Measuring Efficiency by Counting Operations

## Approach #2: Time Complexity Analysis

- ▶ Determine efficiency of code **without** running it.
- ▶ Idea: find a formula for time taken as a function of input size.

# **Advantages of Time Complexity**

1. Doesn't depend on the computer.
2. Reveals which inputs are “hard”, which are “easy”.
3. Tells us how algorithm scales.

## Exercise

Write a function `mean` which takes in a NumPy array of floats and outputs their mean.

```
def mean(numbers):
    total = 0
    n = len(numbers)
    for x in numbers:
        total += x
    return total / n
```

# Time Complexity Analysis

- ▶ How long does it take mean to run on an array of size  $n$ ? Call this  $T(n)$ .
- ▶ We want a formula for  $T(n)$ .
- ▶ **Idea:**
  - ▶ Assume certain basic operations (like adding two numbers) take a constant amount of time.
  - ▶ Count total time taken by basic operations.

# Basic Operations with Arrays

We'll assume that these operations on NumPy arrays take **constant time**.

- ▶ accessing an element: `arr[i]`
- ▶ asking for the length: `len(arr)`

# Example

|  | Time/exec. | # of execs. |
|--|------------|-------------|
| <pre>def mean(numbers):<br/>    total = 0<br/>    n = len(numbers)<br/>    for x in numbers:<br/>        total += x<br/>    return total / n</pre> |            |             |

# Example: mean

- ▶ Total time:

$$\begin{aligned}T(n) &= c_3(n + 1) + c_4n + (c_1 + c_2 + c_3) \\&= (c_3 + c_4)n + (c_1 + c_2 + c_3)\end{aligned}$$

- ▶ “Forgetting” constants, lower-order terms with “Big-Theta”:  $T(n) = \Theta(n)$ .
- ▶  $\Theta(n)$  is the **time complexity** of the algorithm.

## Main Idea

Forgetting constant, lower order terms allows us to focus on how the algorithm **scales**, independent of which computer we run it on.

# Careful!

- ▶ Not always the case that a single line of code takes constant time per execution!

# Example

Time/exec. # of execs.

```
def mean_2(numbers):
    total = sum(numbers)
    n = len(numbers)
    return total / n
```

## Example: mean\_2

- ▶ Total time:

$$T(n) = c_1 n + (c_0 + c_2 + c_3)$$

- ▶ “Forgetting” constants, lower-order terms with “Big-Theta”:  $T(n) = \Theta(n)$ .

## Exercise

Write an algorithm for finding the maximum of an array of  $n$  numbers. What is its time complexity?

Time/exec. # of execs.

```
def maximum(numbers):
    current_max = -float('inf')
    for x in numbers:
        if x > current_max:
            current_max = x
    return current_max
```

## Main Idea

Using Big-Theta allows us not to worry about *exactly* how many times each line runs.