

DSC 40B

Theoretical Foundations II

Average Case

Average Case

- ▶ Best case and worst case can be **misleading**.
 - ▶ Depend on a **single good/bad input**.
- ▶ How long does algorithm take on a **typical input**?

Linear Search

- ▶ **Best Case:** target is right at the beginning, $\Theta(1)$.
- ▶ **Worst Case:** target is at end (or missing), $\Theta(n)$.
- ▶ **Typical Case:** target is somewhere in the middle, $\Theta(n/2) = \Theta(n)$.

More Formally

- ▶ Recall: the **expectation**.

winnings	probability
\$ 0	50%
\$ 1	30%
\$ 10	18%
\$ 50	2%

- ▶ Expected winnings:

Expected Time

- ▶ We'll compute the **expected time**:

$$T_{\text{avg}}(n) = \sum_{\text{case} \in \text{all cases}} P(\text{case}) \cdot T(\text{case})$$

- ▶ Also called the **average case time complexity**.

Average Case Complexity

- ▶ Step 1: Determine the possible cases.
- ▶ Step 2: Determine the probability of each case.
- ▶ Step 3: Determine the time taken for each case.
- ▶ Step 4: Compute the expected time (average).

Step 1: Determine the Cases

- ▶ Example: linear search.
 - Case 1: target is first element
 - Case 2: target is second element
 - :
 - Case n : target is n th element
 - Case $n + 1$: target is not in array

Step 2: Case Probabilities

- ▶ What is the probability that we see each case?
 - ▶ Example: what is the probability that the target is the first element?
- ▶ Often requires making **assumptions**.
 - ▶ Example: target is equally likely to be in any position, but *must* be in array.
- ▶ Often guarantee assumptions by **randomizing**.
 - ▶ Example: randomly shuffle input array.

Example

- ▶ **Assume:** target is in the array exactly once
- ▶ **Randomize** the array (in linear time)
- ▶ Each case is **equally likely**. Probability: $1/n$.

Step 3: Case Times

- ▶ Example: linear search.
 - ▶ Let's say it takes time c per iteration.

Case 1: time c

Case 2: time $2c$

⋮

Case k : time $c \cdot k$

⋮

Case n : time $c \cdot n$

Step 4: Compute Expectation

$$T_{\text{avg}}(n) = \sum_{i=1}^n P(\text{case } i) \cdot T(\text{case } i)$$

Average Case Time Complexity

- ▶ The **average case** time complexity of **linear search** is $\Theta(n)$, as expected.

Note

- ▶ **Worst case** time complexity is still useful.
- ▶ Easier to calculate.
- ▶ Often same as average case (but not always!)
- ▶ Sometimes worst case is very important.
 - ▶ Real time applications, time complexity attacks

DSC 40B

Theoretical Foundations II

Average Case in Movie Problem

The Movie Problem

```
def find_movies(movies, t):
    n = len(movies)
    for i in range(n):
        for j in range(i + 1, n):
            if movies[i] + movies[j] == t:
                return (i, j)
    return None
```

Exercise

What is the best case time complexity of `find_movies`? What about the worst case?

Time Complexity

- ▶ Best case: $\Theta(1)$
- ▶ Worst case: $\Theta(n^2)$
- ▶ Average case: $\Theta(?)$

Step 1: Determine the Cases

- ▶ Each possible pair of movies is a case.
- ▶ There are $\binom{n}{2}$ cases.

Step 2: Case Probabilities

- ▶ **Assume:** there is a *unique* pair that adds to t.
- ▶ **Assume:** all pairs are equally likely.
- ▶ Probability of any case: $\frac{1}{\binom{n}{2}} = \frac{2}{n(n-1)}$

Step 3: Case Time

- ▶ How much time is taken for a particular case?
- ▶ Suppose the 4th and 11th movie sum to the target.
- ▶ How long does it take to find this pair?

```
1 def find_movies(movies, t):
2     n = len(movies)
3     for i in range(n):
4         for j in range(i + 1, n):
5             if movies[i] + movies[j] == t:
6                 return (i, j)
7
8 return None
```

Exercise

Roughly how many times does line 5 execute if the 4th and 11th movies add to the target?

Visualization

	1	2	3	\dots	j	\dots	$n-1$	n
1	m							
2	m	m						
3	m	m	m					
\vdots								
i	m	m	m	m				
\vdots								
$n-1$	m	m	m	m	m			
n	m	m	m	m	m	m		

Average Case

- ▶ The average case time complexity is $\Theta(n^2)$.

DSC 40B

Theoretical Foundations II

Lower Bound Theory

Problems and Algorithms

- ▶ There can be many **algorithms** for solving the same **problem**.
- ▶ Some have better time complexity than others.
- ▶ **Very important question:** For a given problem, what is the **best possible** time complexity?

Lower Bounds

- ▶ No algorithm can have a better (worst case) time complexity than a **theoretical lower bound**.
- ▶ E.g.: $\Omega(n)$ is a lower bound for movie problem.
- ▶ Any algorithm for the movie problem must take $\Omega(n)$ time in the worst case.

Definition

$f(n)$ is a **theoretical lower bound** for a problem if every possible algorithm's worst case complexity is $\Omega(f(n))$.

Main Idea

No algorithm's worst case can be better than theoretical lower bound.

Linear Search

- ▶ **Given:** an array arr of numbers and a target t .
- ▶ **Find:** the index of t in arr , or **None** if it is missing.
- ▶ Theoretical Lower Bound: $\Omega(n)$.
 - ▶ Why? In the worst case, *every* algorithm has to look through all n numbers.

Useless Lower Bounds

- ▶ $\Omega(1)$ is also a theoretical lower bound.
- ▶ But it is useless... no algorithm exists.

Tight Lower Bounds

- ▶ A lower bound is **tight** if there exists an algorithm with that worst case time complexity.
- ▶ That algorithm is in a sense **optimal**.

The Question

- ▶ What is the **best possible** time complexity for an algorithm solving the movie problem?
- ▶ $\Omega(n)$ is a non-trivial lower bound.
 - ▶ Must read every movie duration in worst case.
- ▶ Is this a tight bound? Is there such an algorithm?

DSC 40B

Theoretical Foundations II

Matrix Multiplication

It's Important

- ▶ Matrix multiplication is a *very* common operation in machine learning algorithms.
- ▶ **Estimate:** 75% - 95% of time training a neural network is spent in matrix multiplication.

Recall

- ▶ If A is $m \times p$ and B is $p \times n$, then AB is $m \times n$.
- ▶ The ij entry of AB is

$$(AB)_{ij} = \sum_{k=1}^p a_{ik} b_{kj}$$

Recall

$$(AB)_{ij} = \sum_{k=1}^p a_{ik} b_{kj}$$

$$\begin{pmatrix} 1 & 2 & 3 \\ 4 & 5 & 6 \\ 7 & 8 & 9 \end{pmatrix} \begin{pmatrix} 5 & -1 \\ 1 & 7 \\ -2 & -3 \end{pmatrix} = \begin{pmatrix} \quad & \quad \\ \quad & \quad \\ \quad & \quad \end{pmatrix}$$

Naïve Algorithm

- ▶ This algorithm is relatively straightforward to code up.

```
def mmul(A, B):
    """
    A is (m x p) and B is (p x n)
    """
    m, p = A.shape
    n = B.shape[1]

    C = np.zeros((m, n))

    for i in range(m):
        for j in range(n):
            for k in range(p):
                C[i,j] += A[i,k] * B[k, j]

    return C
```

Time Complexity

- ▶ The naïve algorithm takes time $\Theta(mnp)$.
- ▶ If both matrices are $n \times n$, then $\Theta(n^3)$ time.
- ▶ **Cubic!**

Cubic Time Complexity

- ▶ The largest problem size that can be solved, if a basic operation takes 1 nanosecond.


1 s	10 m	1 hr
1,000	6,694	15,326

The Question

- ▶ Can we do better?
- ▶ How fast can we possibly multiply matrices?


Theoretical Lower Bound

- ▶ If A and B are $n \times n$, C will have n^2 entries.
- ▶ Each entry must be filled: $\Omega(n^2)$ time.
- ▶ That is, matrix multiplication must take at least quadratic time.
- ▶ Is this bound **tight**? Can it be increased?



Strassen's Algorithm


- ▶ Cubic was as good as it got...
- ▶ ...until Strassen, 1969.
- ▶ Time complexity: $\Theta(n^{\log_2 7}) = \Theta(n^{2.8073})$



Currently


- ▶ The fastest¹ known matrix multiplication algorithm is due to Le Gall.
- ▶ $\Theta(n^{2.3728639})$ time.

¹In terms of asymptotic time complexity.



Interestingly...

- ▶ No one knows what the lowest possible time complexity is.
- ▶ It could be $\Theta(n^2)$!
- ▶ The “best” matrix multiplication algorithm is probably still undiscovered.



Irony

- ▶ There are many matrix multiplication algorithms.
- ▶ How fast is numpy's matrix multiply?
- ▶ $\Theta(n^3)$.

Why?

- ▶ Strassen *et al.* have better asymptotic complexity.
- ▶ But much (much!) larger “hidden constants”.
- ▶ Remember, which is better for small n :
 $999,999n^2$ or n^3 ?

Optimization

- ▶ Numpy, most others use **highly optimized** cubic time algorithms.

Main Idea

No one knows what the lowest possible time complexity of matrix multiplication is, and some algorithms are approaching $\Theta(n^2)$.

But most useful implementations take $\Theta(n^3)$ time.