

DSC40B: Theoretical Foundations of Data Science II

Lecture 8: *Binary search tree*

Instructor: Yusu Wang

Set operations

- ▶ Imagine you are maintaining a database indexed by some keys (real values), and you hope to support the following operations:

- ▶ Search
- ▶ Maximum
- ▶ Minimum
- ▶ Successor
- ▶ Predecessor

First approach: sort the array of keys

$\Theta(\lg n)$
 $\Theta(1)$
 $\Theta(1)$
 $\Theta(1)$
 $\Theta(1)$

- ▶ Insert
- ▶ Delete
- ▶ Extract-Max
- ▶ Increase-key

Can we use a linked list instead?

$\Theta(n)$

How to have a good data structure so we can support all these operations efficiently?

Today

- ▶ **Binary search tree**
 - ▶ support all the operations from previous slide
 - ▶ in time proportional to height of tree
- ▶ (Review): how to implement key operations, and time complexity
 - ▶ search, insert (and delete)
- ▶ Extension to **balanced** binary search tree
- ▶ *Select query: augmenting data structure*
 - ▶ median, order statistics




What is binary search tree?




Binary tree

- ▶ A binary tree is a rooted tree
 - ▶ where each node has at most 2 children
- ▶ Represented by a linked data structure
- ▶ Each node contains at least fields:
 - ▶ *Key*
 - ▶ *Left*
 - ▶ *Right*
 - ▶ *Parent*




Example

- ▶ From root, following left pointers, we will visit
 - ▶ 13, 6, 3, 2, *Nil*



Binary tree

- ▶ A binary tree is a rooted tree where
 - ▶ each node has at most 2 children
- ▶ A node is the root of the tree
 - ▶ if its parent is Nil
- ▶ A node is a leaf
 - ▶ if both children are Nil
- ▶ Left sub-tree, right sub-tree
- ▶ A complete binary tree
 - ▶ is a binary tree where each node has two children other than leaves



Binary search tree (BST)

- ▶ **Binary-search-tree property**
 - ▶ For any node $x \in T$,
 $x.Key \geq (x.Left).Key$ and $x.Key \leq (x.Right).Key$
- ▶ A binary tree T is a **binary search tree (BST)** if
 - ▶ it satisfies the binary search tree property

Fact

Given a binary search tree T , for any node $x \in T$, we have

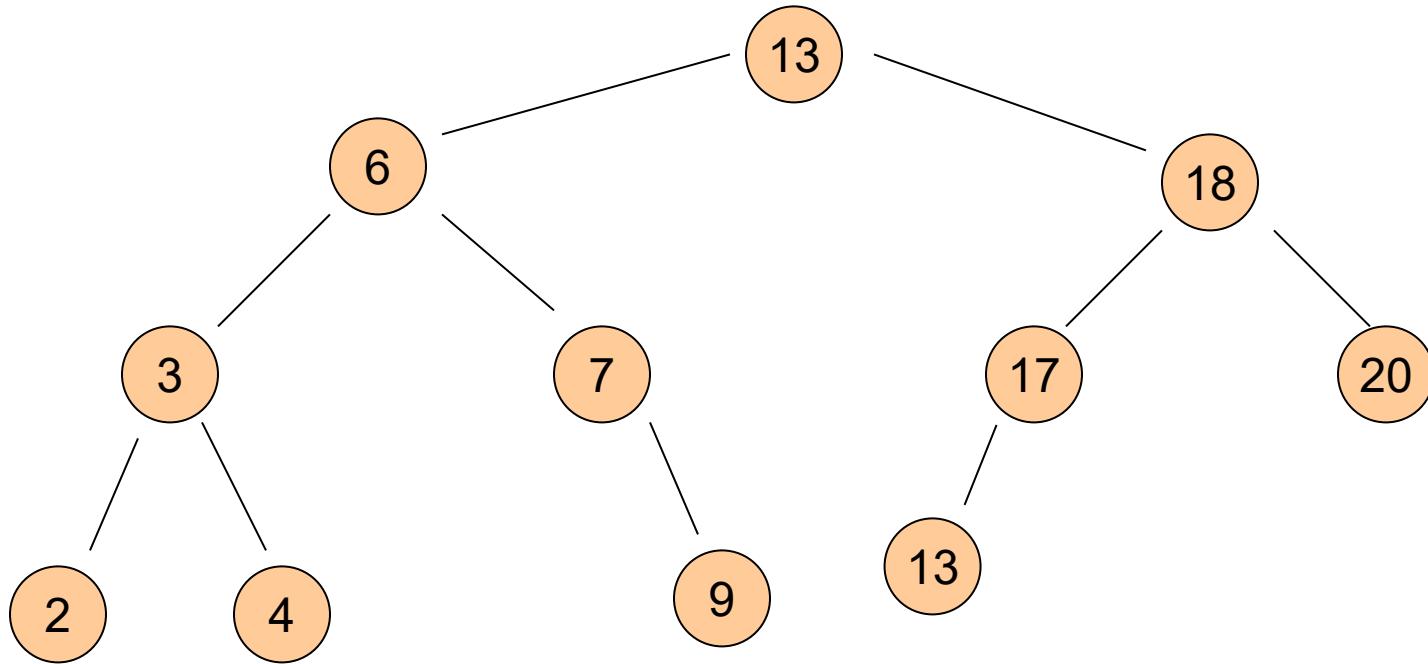
$x.Key \geq y.Key$ if y is in the left subtree of x ; and

$x.Key \leq y.Key$ if y is in the right subtree of x




Example

► A valid BST



Example

► ?

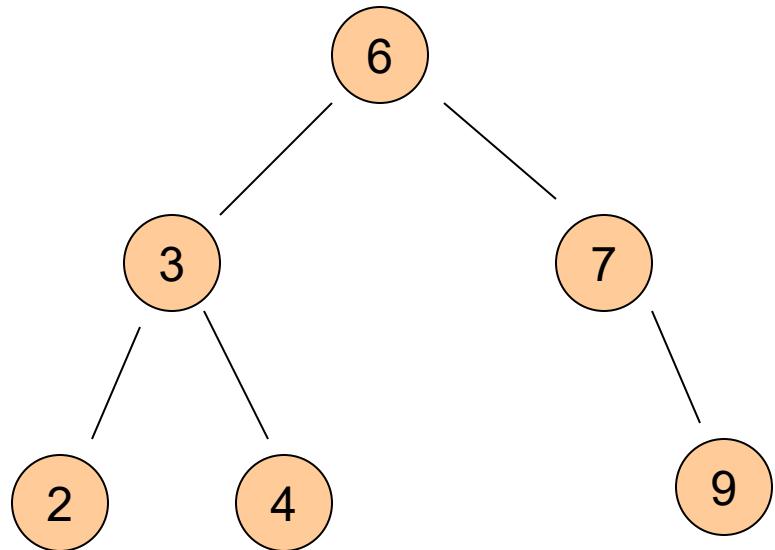


►

Properties

- ▶ Given the same set of elements
 - ▶ there are many possible BSTs over them

- ▶ Minimum?
 - ▶ Does it have to be a leaf?



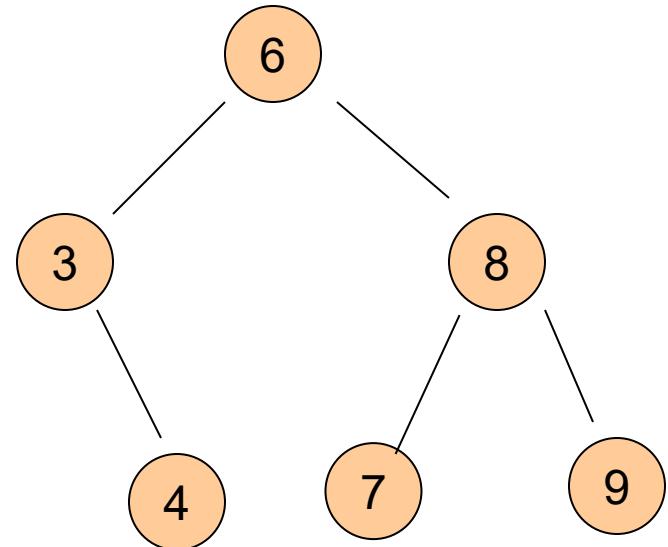
Properties

- ▶ Given the same set of elements
 - ▶ there are many possible BSTs over them

- ▶ Minimum?
 - ▶ Does it have to be a leaf?

- ▶ Maximum?

- ▶ Given n nodes,
 - ▶ Tallest possible BST tree has height $h = \underline{n}$
 - ▶ Shortest possible BST tree has height $h = \underline{\log_2 n} = \Theta(\lg n)$



Operations in BST




Search operation

- ▶ A BST T with n nodes can be viewed as a way to store n keys in a smart way, so that queries among these keys become easy.
- ▶ Tree-search(x, k)
 - ▶ Input: given a tree node x and a query key k
 - ▶ Output: search whether k is in the tree rooted at x
 - ▶ if it is in, return a node y s.t. $y.key = k$
 - ▶ otherwise, returns NIL

Tree-search($x, 8$)

Tree-search($x, 4$)

Tree-search($x, 5$)



Tree-search algorithm, recursive version

Tree-search (x, k)

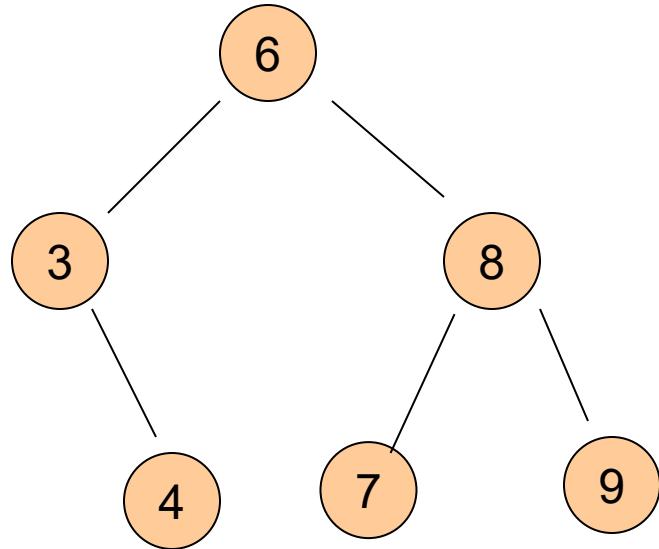
if $x = \text{Nil}$ or $k = x.\text{key}$

then return x

if $k < x.\text{key}$

then return Tree-search($x.\text{left}, k$)

else return Tree-search($x.\text{right}, k$)



- ▶ Given an input tree T and a key k
 - ▶ we will start by calling Tree-search($T.\text{root}, k$)



Tree-search algorithm, recursive version

Tree-search (x, k)

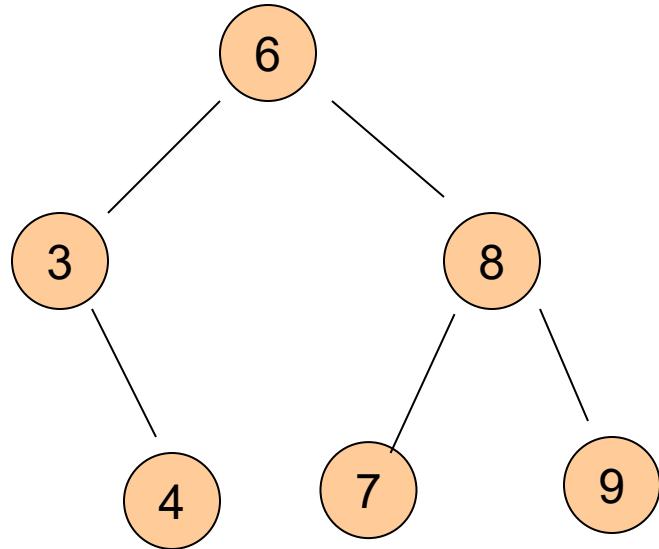
if $x = \text{Nil}$ or $k = x.\text{key}$

then return x

if $k < x.\text{key}$

then return Tree-search($x.\text{left}, k$)

else return Tree-search($x.\text{right}, k$)



- ▶ Time complexity analysis

- ▶ let $T(n)$ denote the worst case time complexity of procedure Tree-search() on any tree of n nodes



Tree-search algorithm, recursive version

Tree-search (x, k)

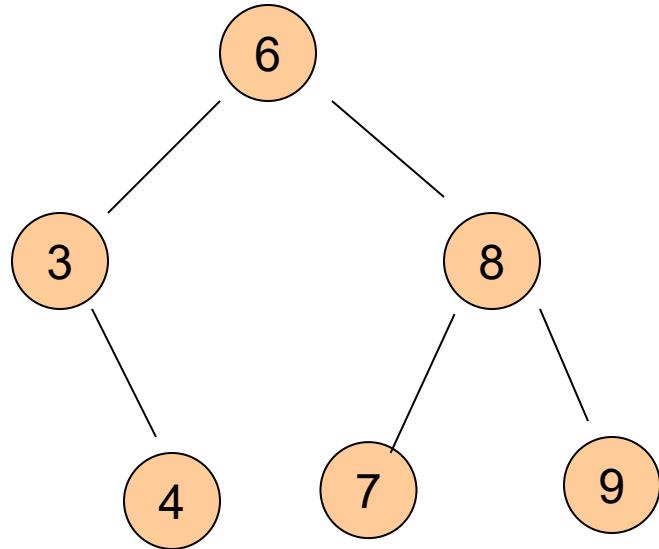
if $x = \text{Nil}$ or $k = x.\text{key}$

then return x

if $k < x.\text{key}$

then return Tree-search($x.\text{left}, k$)

else return Tree-search($x.\text{right}, k$)



► Time complexity analysis

- other than recursive call, $\Theta(1)$ within each Tree-search call
- thus, $T(n)$ is proportional to the number of nodes x we will call Tree-search on
- $T(n) = \Theta(\text{tree-height}) = O(n)$



Tree-search: iterative version

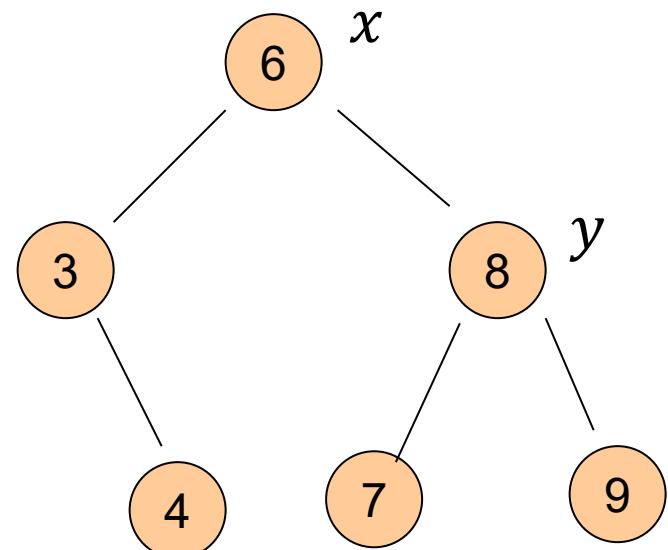
```
procedure IterativeTreeSearch(x,K)
1 while (x = NIL) and (K ≠ x.key) do
2     if (K ≤ x.key) then
3         |   x ← x.left;
4     else
5         |   x ← x.right;
6     end
7 end
8 return (x);
```



Minimum / Maximum

▶ Tree-minimum(x)

- ▶ **Input:** a node x of a BST T
- ▶ **Output:** return the node containing minimum key in the subtree rooted at x



Minimum / Maximum

▶ Tree-minimum(x)

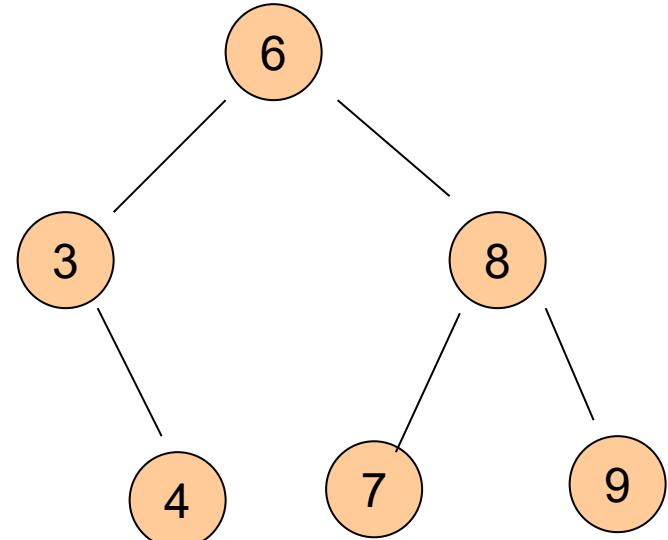
- ▶ **Input:** a node x of a BST T
- ▶ **Output:** return the node containing a minimum key in the subtree rooted at x

```
Tree-minimum( $x$ )
```

```
    while ( $x.left \neq Nil$ )
        do  $x = x.left$ ;
    return  $x$ ;
```

▶ Time complexity

- ▶ $T(n) = \Theta(h)$ where h is height of input tree



Minimum / Maximum

▶ Tree-maximum(x)

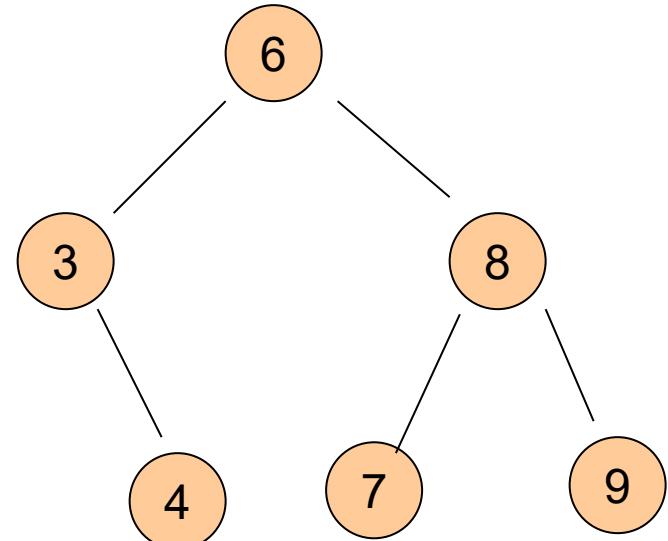
- ▶ **Input:** a node x of a BST T
- ▶ **Output:** return the node containing a maximum key in the subtree rooted at x

```
Tree-maximum( $x$ )
```

```
    while (  $x.right \neq Nil$  )  
        do  $x = x.right$ ;  
    return  $x$ ;
```

▶ Time complexity

- ▶ $T(n) = \Theta(h)$ where h is height of input tree




Tree-insert

- ▶ **Tree-insert(x, k)**
 - ▶ Input: a BST tree node x and a key k
 - ▶ Output: insert k to the tree rooted at x such that the resulting tree is still a binary search tree



Examples



Tree-Insert(T.root, 8)

Tree-Insert(T.root, 6.5)

Use tree-search !

Tree-insert

Tree-insert(T, k)

$y = \text{Nil}; x = T.\text{root}$

$z.\text{key} = k; z.\text{left} = \text{Nil}; z.\text{right} = \text{Nil}$

while ($x \neq \text{Nil}$) do

$y = x$

if ($z.\text{key} < x.\text{key}$)

then $x = x.\text{left}$

else $x = x.\text{right}$

$z.\text{parent} = y$

if ($y = \text{Nil}$) then $T.\text{root} = z$

else if ($z.\text{key} < y.\text{key}$)

then $y.\text{left} = z$

else $y.\text{right} = z$

- z is the new node to be inserted
- Locate potential parent y of z .

- Set up z as appropriate child of y



Tree-insert

Tree-insert(T, k)

$y = \text{Nil}; x = T.\text{root}$

$z.\text{key} = k; z.\text{left} = \text{Nil}; z.\text{right} = \text{Nil}$

while ($x \neq \text{Nil}$) do

$y = x$

 if ($z.\text{key} < x.\text{key}$)

 then $x = x.\text{left}$

 else $x = x.\text{right}$

$z.\text{parent} = y$

 if ($y = \text{Nil}$) then $T.\text{root} = z$

 else if ($z.\text{key} < y.\text{key}$)

 then $y.\text{left} = z$

 else $y.\text{right} = z$

▶ Time complexity

▶ $T(n) = \Theta(h)$, where h is height of input tree

Summary

- ▶ Suppose n input keys are already stored in a BST of height h

- ▶ Search
- ▶ Maximum
- ▶ Minimum
- ▶ Successor
- ▶ Predecessor

- ▶ Insert
- ▶ Delete
- ▶ Extract-Max
- ▶ Increase-key

Time complexity

$\Theta(h)$

$\Theta(h)$

$\Theta(h)$

$\Theta(h)$

$\Theta(h)$

$\Theta(h)$

$\Theta(h)$

$\Theta(h)$

$\Theta(h)$

- However, performance depending on height!
- Height $h = O(n)$ and $h = \Omega(\lg n)$

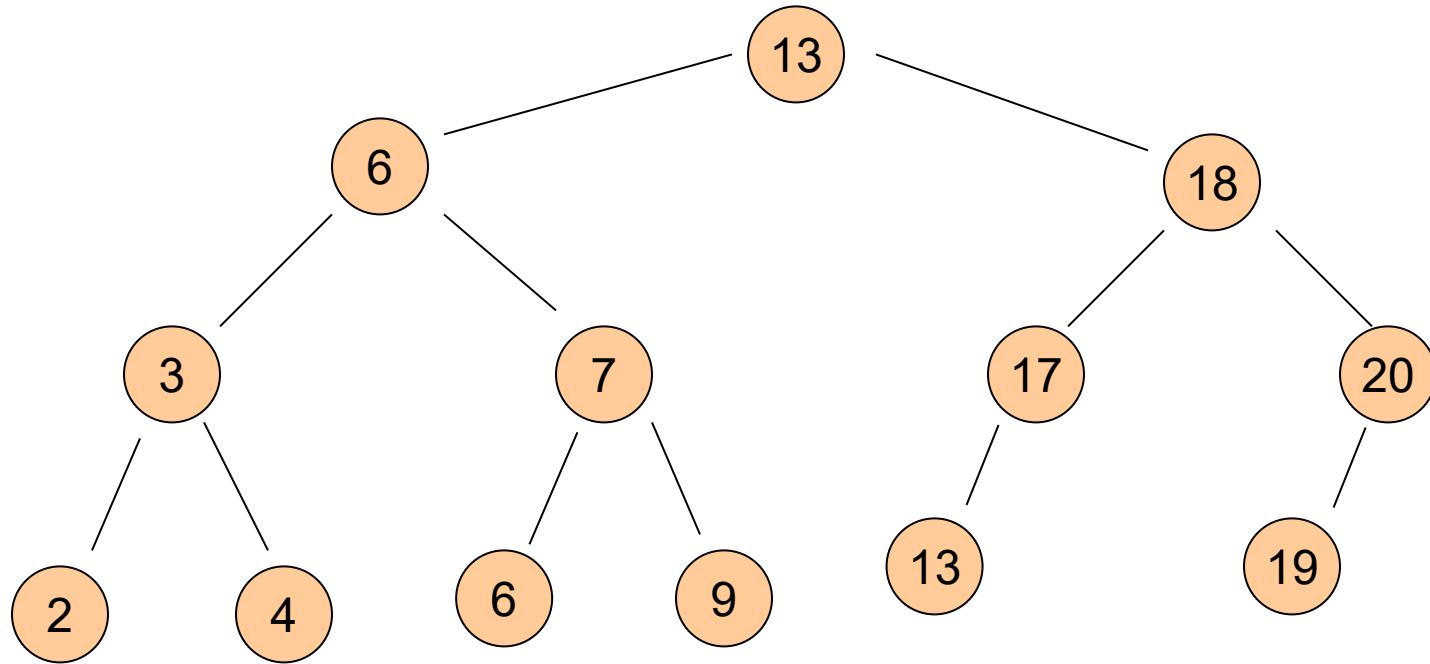
- To have good performance, we want to keep the tree height low!



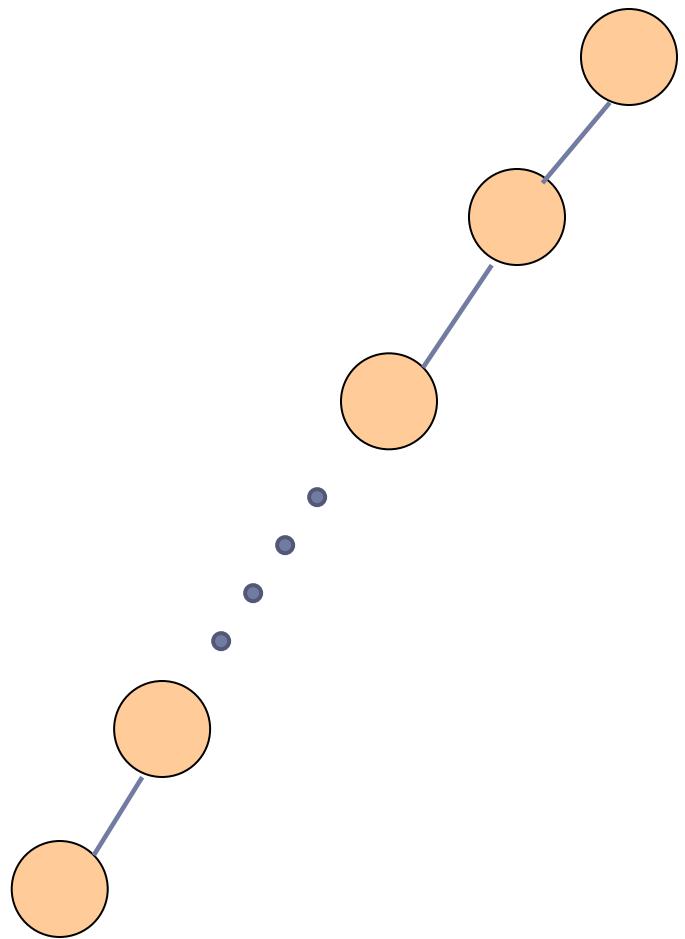
Balanced binary search tree



Good tree



Bad Tree




Balanced binary search tree

- ▶ It turns out that there are ways to add extra conditions to binary search trees, so that their height is $\Theta(\lg n)$
 - ▶ E.g, red-black tree, AVL tree, etc
- ▶ Once such a tree is created,
 - ▶ it can support search, minimum, maximum etc in $\Theta(h) = \Theta(\lg n)$ time using the same algorithms described before
 - ▶ the extra work comes at handling dynamic operations: insertion, deletion, and so on. Re-balancing is needed
 - ▶ however, for standard balanced BSTs, all these operations can be handled in $\Theta(\lg n)$ time.



Rotation operation

- ▶ Left rotation or Right rotation to keep tree height low



With balanced BST

- ▶ Suppose n input keys are already stored in a balanced BST

	Time complexity
▶ Search	$\Theta(\lg n)$
▶ Maximum	$\Theta(\lg n)$
▶ Minimum	$\Theta(\lg n)$
▶ Successor	$\Theta(\lg n)$
▶ Predecessor	$\Theta(\lg n)$
▶ Insert	$\Theta(\lg n)$
▶ Delete	$\Theta(\lg n)$
▶ Extract-Max	$\Theta(\lg n)$
▶ Increase-key	$\Theta(\lg n)$

- Height of tree will be $\Theta(\lg n)$, where n is number of nodes in the tree



Select queries
augmenting data structure



-
- ▶ What if we also want to perform Select operation
 - ▶ Select (T, k):
 - ▶ Given a list of records whose keys are stored in T , return the node whose key has rank k .
 - ▶ We can do linear search to find it. But can we do better?
 - ▶ Goal:
 - ▶ **Augment** the binary search tree data structure so as to support Select (T, k) efficiently
-

In particular,

- ▶ Select (T, k)
- ▶ Goal:
 - ▶ Augment the binary search tree data structure so as to support Select (T, k) efficiently
- ▶ Ordinary binary search tree
 - ▶ $O(h)$ time for Select(T, k)
- ▶ Using balanced search tree)
 - ▶ $O(\lg n)$ time for Select(T, k)




How do we augment a BST T?

- ▶ At each node x of the tree T
 - ▶ store $x.size = \# \text{ nodes in the subtree rooted at } x$
 - ▶ Include x itself
 - ▶ If a node (leaf) is NIL, its size is 0.
- ▶ Space of an augmented tree:
 - ▶ $\Theta(n)$
- ▶ Basic property:
 - ▶ $x.size = x.left.size + x.right.size + 1$



An example



How to setup size information?

▶ procedure *AugmentSize*(*treenode* *x*)

If (*x* ≠ *NIL*) then

Lsize = *AugmentSize*(*x.left*);

Rsize = *AugmentSize*(*x.right*);

x.size = *Lsize* + *Rsize* + 1;

 Return(*x.size*);


end

Return (0);

Postorder traversal
of the tree !



How to perform select with aug-BST?



-
- ▶ Let T be an augmented binary search tree
 - ▶ $BST\text{-}Select(x, k)$:
 - ▶ Return the k -th smallest element in the subtree rooted at x
 - ▶ $BST\text{-}Select(T.\text{root}, k)$ returns the k -th smallest elements in the entire tree.



▶ procedure *BST-Select*(*treenode* x , k)
 $\ell = x.\text{left.size} + 1$
 If ($k == \ell$) **then**
 Return x
 elseif ($k < \ell$)
 Return *BST-Select*($x.\text{left}$, k)
 else Return *BST-Select*($x.\text{right}$, $k - \ell$)

- ▶ Initial call: *BST-Select*($T.\text{root}$, k)
- ▶ Time complexity:
- ▶ $T(n) = \Theta(\text{height of tree}) = O(n)$
 - ▶ If T is balanced, then $T(n) = \Theta(\lg n)$

Are we done?

- ▶ Need to maintain the augmented information under dynamic changes of the tree!
 - ▶ i.e, under insertions / deletions
 - ▶ in this case, just adjusting this size count as we update nodes, or under rotations, and it does not increase asymptotic time complexity of these operations
- ▶ Remark:
 - ▶ Select() in an sorted array can be done in $\Theta(1)$ time.
 - ▶ However, an array does not support dynamic operations (insert/delete) efficiently. That's augmented BST is a better data structure in this case.



Summary

- ▶ Simple example of augmenting data structures
- ▶ In general, the augmented information can be quite complicated
 - ▶ Can be a separate data structure!
- ▶ Need to consider how to maintain such information under dynamic changes



FIN

