

DSC40B: Theoretical Foundations of Data Science II

Lecture 9: *Hashing and Hash table*

Instructor: Yusu Wang

Recall: Set operations

- ▶ Imagine you are maintaining a database indexed by some keys (real values), and you hope to support the following operations:
 - ▶ Search
 - ▶ Maximum
 - ▶ Minimum
 - ▶ Successor
 - ▶ Predecessor
- ▶ Insert
- ▶ Delete
- ▶ Extract-Max
- ▶ Increase-key

Using balanced BST, all these operations can be done in $O(\lg n)$ time



Dictionary operations

- ▶ Given a universe of elements U
 - ▶ these elements may not be numbers
- ▶ Need to store some keys
- ▶ Need to perform the following operations for keys
 - ▶ Insert
 - ▶ Search
 - ▶ Delete

We can use balanced BST for this purpose if keys can be compared.

But we can have even lighter weight data structure to handle these



Today

- ▶ Hashing in general
- ▶ Hash table
 - ▶ For dictionary operations



Hashing



Hashing

- ▶ An important idea used commonly in practice
- ▶ Many uses:
 - ▶ Fast queries on a large data set.
 - ▶ Verifying message integrity.
 - ▶ Identify if file has changed in version control.



Hash function

- ▶ Mathematically, a hash function is simply a function
 - ▶ $f: U \rightarrow X$ from one set to another
- ▶ To make it useful, in practice,
 - ▶ It often maps some potentially large or complex object to a much smaller and simpler “fingerprint” or “signature”
 - ▶ One also want to mapping to be “uniform” and cause few “collisions”.
 - ▶ Note: a hash function needs to be deterministic!
 - ▶ Hashing the same object twice, we should get the same answer



Some examples

► A cryptographic hash function:

- ▶ maps data of arbitrary size into an often fixed sized output of much smaller size
- ▶ e.g, MD5: maps it to a 128-bit value
- ▶ hard to “reverse engineer” input from hash.
- ▶ two similar input could and should lead to very different hash values

```
▶ > echo "My name is Justin" | md5  
▶ a741d8524a853cf83ca21eabf8cea190  
▶ > echo "My name is Justin!" | md5  
▶ f11eed2391bbd0a5a2355397c089fafd
```

```
▶ > md5 slides.pdf  
▶ e3fd4370fda30ceb978390004e07b9df
```

A cryptographic hash function

- ▶ Why?
 - ▶ I release a piece of software.
 - ▶ I host it on Google Drive.
 - ▶ Someone (Google, US Gov., etc.) decides to insert extra code into software to spy on users.
 - ▶ You have no way of knowing.
- ▶ What do I do?
 - ▶ I release a piece of software and **publish the hash**
 - ▶ I host it on Google Drive.
 - ▶ Someone inserts extra code into software to spy on users.
 - ▶ You download the software and hash it. If hash is different, you know the file has been changed!



Some examples

- ▶ Want to place images into 100 bins.
 - ▶ How do we decide which bin an image goes into?
 - ▶ Hash function!
 - ▶ Takes in an image.
 - ▶ Outputs a number in $\{1, 2, \dots, 100\}$



Hash Table



Dictionary operations

- ▶ Given a universe of elements U
- ▶ Need to store some keys
- ▶ Need to perform the following operations for keys
 - ▶ Insert
 - ▶ Search
 - ▶ Delete
- ▶ In other words, the key operation is **membership queries** (search), but also allow dynamic updates (insert, delete).



Approach 0

- ▶ Use an array to organize all the keys
 - ▶ We could pre-sort the array.
 - ▶ Preprocessing time:
 - ▶ Search:
 - ▶ Insert / Delete:



Approach 1

- ▶ Organize all keys in a doubly-linked list
 - ▶ Pre-processing: **None**
 - ▶ Insert:
 - ▶ Delete
 - ▶ Search:
 - ▶ Space:



Approach 2

- ▶ Organize all keys in a balanced BST
 - ▶ Search:
 - ▶ Insert:
 - ▶ Delete
- ▶ Space:



Approach 3

- ▶ Direct address tables
 - ▶ Suppose we know that all the keys we ever care are from 0 to $N = 99,999$
 - ▶ e.g, we are querying for zipcodes
 - ▶ Open an array A of size N
 - ▶ If a zipcode z is in, set $A[z] = 1$, and 0 otherwise
 - ▶ Given a query zipcode, say 223300, simply return $A[223300]$
 - ▶ Search:
 - ▶ Insert:
 - ▶ Delete:
- What's the problem with this approach?
- Not practical for if the size of universe U (which is N in this example) is huge!



Use Hash Table!

- ▶ U : universe
- ▶ $T[0 \dots m - 1]$: a hash table of size m
 - ▶ $m \ll |U|$
 - ▶ usually, we choose m to be around the size of data we will see
- ▶ Hash functions
 - ▶ $h: U \rightarrow \{0, 1, \dots, m - 1\}$
- ▶ $h(k)$ is called the **hash value** of key k .
 - ▶ Given a key k , we will store it in location $h(k)$ of hash table T ,
 - ▶ i.e, store it at $T[h(k)]$

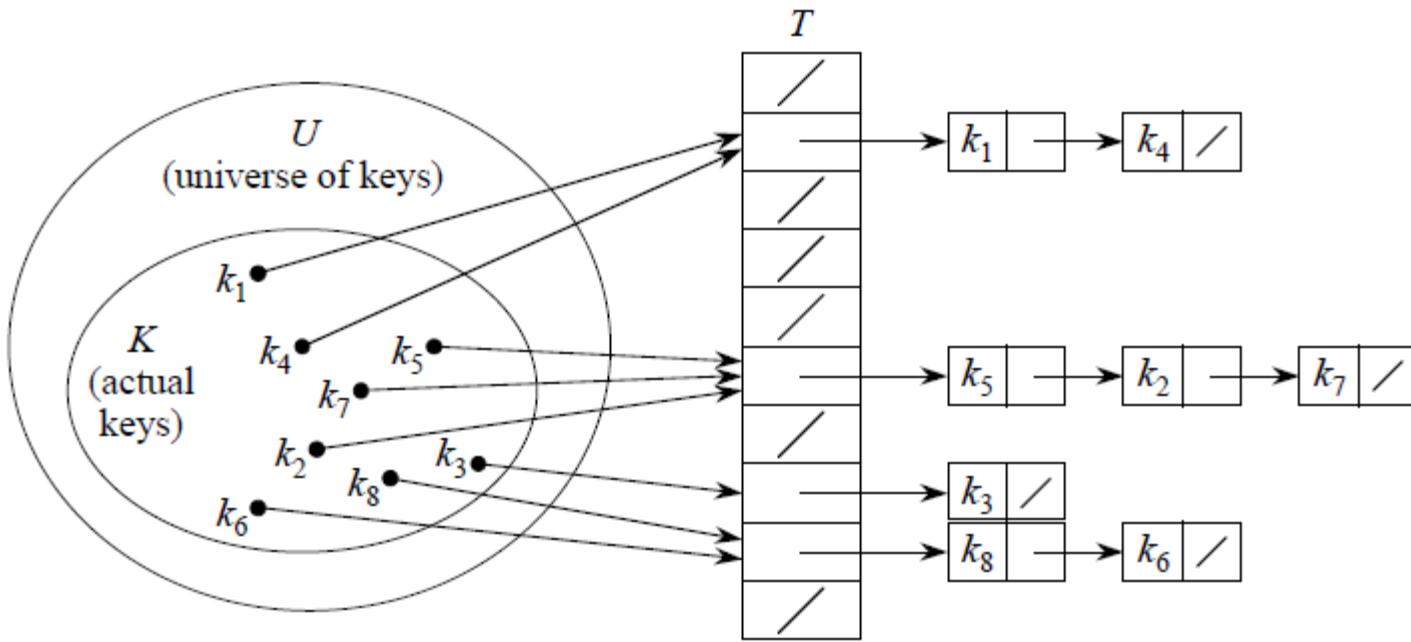


Collision

- ▶ Since the size of hash table is smaller than the universe:
 - ▶ Multiple keys may hash to the same slot.
 - ▶ A **collision** happens when $h(x) = h(y)$ for $x \neq y \in U$
- ▶ How to handle collisions?
 - ▶ Chaining
 - ▶ Open addressing
 - ▶ ...



Collision Resolved by Chaining



- ▶ $T[j]$: a pointer to the head of the linked list of all stored elements that hash to j
- ▶ **Nil** otherwise

A simple example

- ▶ U : all positive integers \mathbb{N}
- ▶ Hash table $T[0, \dots, 10]$
- ▶ Hash function: $h: \mathbb{N} \rightarrow [0, \dots, 10]$
 - ▶ $h(x) = x \bmod 11$



Dictionary operations

- ▶ **Chained-Hash-Insert (T, x)** O(1)
- ▶ Insert x at the head of list $T[h(\text{key}(x))]$
- ▶ **Chained-Hash-Search(T, k)** O(\text{length}(T[h(k)]))
- ▶ Search for an element with key k in list $T[h(k)]$
- ▶ **Chained-Hash-Delete(T, x)** O(\text{length}(T[h(\text{key}(x))]))
- ▶ Delete x from the list $T[h(\text{key}(x))]$



Good Hash Function

- ▶ **Performance of Hash table operations**
 - ▶ depend on the number of elements hashed to each slot in hash table.
- ▶ **Intuitively,**
 - ▶ A good hash function should spread elements into the hash table uniformly
 - ▶ If there are n elements in the input data and m slots
 - ▶ then ideally there should be $\frac{n}{m}$ elements in each hash table slot.



Average case analysis

- ▶ n : # elements in the table
- ▶ m : size of table (# slots in the table)
- ▶ Load factor:
 - ▶ $\alpha = \frac{n}{m}$: average number of elements per linked list
 - ▶ Intuitively the optimal time needed
- ▶ Individual operation can be slow ($O(n)$ time)
 - ▶ *Under certain assumption of the distribution of keys*, analyze expected performance.



Simple uniform hashing assumption

- ▶ Simple uniform hashing assumption:
 - ▶ any given element is equally likely to hash into any of the m slots in T
- ▶ Let n_j be length of list $T[j]$
 - ▶ $n = n_0 + n_1 + \dots + n_{m-1}$
 - ▶ Under simple uniform hashing assumption:
 - ▶ expected value $E[n_j] = \alpha = \frac{n}{m}$

Why?



Why

- ▶ Let $\{k_1, k_2, \dots, k_n\}$ be the set of keys
- ▶ Goal: Estimate $E(n_j)$
- ▶ Let $X_i = \begin{cases} 1 & \text{if } h(k_i) = j \\ 0 & \text{otherwise} \end{cases}$
- ▶ Note: $n_j = \sum_{i=1}^n X_i$!
- ▶ $E[X_i] = \Pr[h(k_i) = j] = \frac{1}{m}$
- ▶ Hence $E[n_j] = E[\sum_{i=1}^n X_i] = \sum_{i=1}^n E[X_i] = \sum_{i=1}^n \frac{1}{m} = \frac{n}{m}$



Expected running time

- ▶ Under Simple uniform hashing assumption

- ▶ Search:

- ▶ Expected time: $ET(n) = \Theta(1 + \frac{n}{m})$
 - ▶ (worst case time $T(n) = \Theta(n)$)

- ▶ Insert:

- ▶ $T(n) = \Theta(1)$

- ▶ Delete:

- ▶ Expected time: $ET(n) = \Theta(1 + \frac{n}{m})$
 - ▶ (worst case time $T(n) = \Theta(n)$)

Key message:

if the load factor $\alpha = \frac{n}{m} = \Theta(1)$

then $\Theta(1)$ expected time for these operations!



Hashing in Python

- ▶ `dict` and `set` implement hash tables

- ▶ Querying is done using `in`:

```
>>> # make a set
```

```
>>> L = {3, 6, -2, 1, 7, 12}
```

```
>>> 1 in L # Theta(1)
```

False

```
>>> 7 in L # Theta(1)
```

True



Some examples of using hash tables



Recall the movie problem

▶ The Movie problem

- ▶ Input: Given a list of length of movies available, stored in array *movies*, and a flight duration D
- ▶ Output: Return two movies whose total length = D ; **None** otherwise.



Recall

- ▶ The naïve algorithm solves it in $\Theta(n^2)$ time
- ▶ But earlier, we mentioned we can do better by
 - ▶ First sorting the array of movie lengths
 - ▶ Then check for each movie x , whether there exists one whose length is $D - \text{length}(x)$
 - ▶ Worst case time complexity $T(n) = \Theta(n \lg n)$



New approach via Hashing

- ▶ Use Hash table, and frame this as a membership query problem

```
def optimize_entertainment_hash(times, D):  
    hash_table = dict()  
    for i, time in enumerate(times):  
        hash_table[time] = i  
  
    for i, time in enumerate(times):  
        target = D - time  
        if target in hash_table:  
            return i, hash_table[target]  
  
    return None
```

Expected time?
 $\Theta(n)$



Another example

- ▶ Anagrams
 - ▶ Two strings w_1 and w_2 are **anagrams** if the letters of w_1 can be permuted to make w_2 .
 - ▶ E.g., “eat” and “tea”, or “listen” and “silent”
- ▶ The *Anagrams problem*
 - ▶ Given a collection of n strings, determine if any two of them are anagrams.



Using Hash table

▶ Observation:

- ▶ two strings are anagrams if their sorted lists are equal
 - ▶ `sorted(w_1) == sorted(w_2)`

```
def any_anagrams(words):  
    seen = set()  
    for word in words:  
        w = sorted(word)  
        if w in seen:  
            return True  
        else:  
            seen.add(w)  
    return False
```

Expected time?
 $\Theta(n)$

Hashing Downsides

- ▶ Only support dictionary queries
 - ▶ i.e, membership queries + insert / delete
 - ▶ Example 1: cannot be used to query for the two movies whose total time is **closest** to D
 - ▶ Example 2: cannot be used for performing a range query in a list of numbers
 - ▶ Say report the number of numbers fall within range $[a, b]$



Hashing Downsides

- ▶ Only support dictionary queries
 - ▶ i.e, membership queries + insert / delete
- ▶ No locality: similar items map to very different bins
 - ▶ Necessarily so for the performance of hashing in most cases!
 - ▶ But, in practice, we often query similar objects continuously
 - ▶ May result in many cache misses, slow



Summary

- ▶ **Hashing**
 - ▶ Very useful idea
 - ▶ Generates a small signature for a potentially large complex object
- ▶ **Hash Table**
 - ▶ Very practical data structure for **dictionary operations**
 - ▶ Very efficient for these operations, can be constant expected time if the load factor $\alpha = \frac{n}{m} = \Theta(1)$
 - ▶ Especially when the number of keys necessary is much smaller than the size of universe where input objects could come from!
 - ▶ In practice, need to choose hash functions properly
 - ▶ There exist intelligent hashing schemes



FIN

