


# DSC 40B

Theoretical Foundations II

## Minimum Spanning Trees

# Today's Problem




- ▶ Choose a set of dirt roads to pave so that:
  - ▶ can get between any two buildings only on paved roads;
  - ▶ total cost is minimized.
- ▶ Solution: compute a **minimum spanning tree**.

# Trees

An undirected graph  $T = (V, E)$  is a **tree** if

- ▶ it is connected; and
- ▶ it is acyclic.

Example: a **tree**. Example: a **tree**. Example: **not** a tree. Example: **not** a tree.




# Trees: Equivalent Definition


An undirected graph  $T = (V, E)$  is a **tree** if

- ▶ it is connected; and
- ▶  $|E| = |V| - 1$ .

Example: a **tree**. Example: a **tree**. Example: **not** a tree. Example: **not** a tree.




# Tree Properties



- ▶ There is a unique simple path between any two nodes in a tree.
- ▶ Adding a new edge to a tree creates a cycle (no longer a tree).
- ▶ Removing an edge from a tree disconnects it (no longer a tree).



# Spanning Trees

Let  $G = (V, E)$  be a **connected** graph. A **spanning tree** of  $G$  is a tree  $T = (V, E_T)$  with the same nodes as  $G$ , and a subset of  $G$ 's edges.




# Many Spanning Trees

The same graph can have many spanning trees.



# Spanning Tree Cost


If  $G = (V, E, \omega)$  is a weighted undirected graph, the **cost** (or **weight**) of a spanning tree is the total weight of the edges in the spanning tree.



Cost:

# Spanning Tree Cost

Different spanning trees of the same graph can have different costs.



Cost:


# Minimum Spanning Tree

- ▶ The **minimum spanning tree** problem is as follows:
  - ▶ GIVEN: A weighted, undirected graph  $G = (V, E, \omega)$ .
  - ▶ COMPUTE: a spanning tree of  $G$  with minimum cost (i.e., minimum total edge weight).
- ▶ For a given graph, the MST may not be unique.

### Exercise

Suppose the edges of a graph  $G = (V, e, \omega)$  all have the same weight. How can we compute an MST of the graph?

# Today's Problem



- ▶ Choose a set of dirt roads to pave so that:
  - ▶ can get between any two buildings only on paved roads;
  - ▶ total cost is minimized.
- ▶ Solution: compute a **minimum spanning tree**.

# MSTs in Data Science?

- ▶ Do we need to find MSTs in data science?
- ▶ Actually, yes! (Next lecture)

# DSC 40B

Theoretical Foundations II

Prim's Algorithm


# Building MSTs

- ▶ How do we build a MST efficiently?
- ▶ We'll adopt a **greedy** approach.
  - ▶ Build a tree edge-by-edge.
  - ▶ At every step, doing what looks best at the moment.
- ▶ This strategy isn't guaranteed to work in all of life's situations, but it works for building MSTs.

# Two Greedy Approaches

- ▶ We'll look at two greedy algorithms:
  - ▶ Today: Prim's Algorithm
  - ▶ Next time: Kruskal's Algorithm
- ▶ Differ in the order in which edges are added to tree.
- ▶ Also differ in time complexity.

# Prim's Algorithm, Informally




- ▶ Start by picking any node to add to “tree”,  $T$ .
- ▶ While  $T$  is not a spanning tree, greedily add **lightest** edge from a node in  $T$  to a node not in  $T$ .
  - ▶ “lightest” = edge of smallest weight
- ▶ **Is this guaranteed to work?** Yes, as we'll see.

# Prim's Algorithm, Equivalently

- ▶ For each node  $u$ , store:
  - ▶ estimated cost of adding node to tree;
  - ▶ estimated “predecessor”  $v$  in the tree.
- ▶ At each step,
  - ▶ Find node with smallest estimated cost.
  - ▶ Add to tree  $T$  by including edge with estimated “predecessor”.
  - ▶ Update cost of neighbors.
- ▶ Same as adding lightest edge from  $T$  to outside  $T$  at every step!

# Prim's Algorithm, Equivalently



- ▶ While  $T$  is not a tree:
  - ▶ find the node  $u \notin T$  with smallest cost
  - ▶ add the edge between  $u$  and its estimated “predecessor” to  $T$
  - ▶ update estimated cost/pred. of  $u$ 's neighbors which aren't already in tree.

# Recall: Priority Queues

- ▶ How do we efficiently find node with smallest cost?
- ▶ Priority Queues:
  - ▶ `PriorityQueue(priorities)`: creates priority queue from dictionary whose values are priorities.
  - ▶ `.extract_min()`: removes and returns key with smallest value.
  - ▶ `.decrease_priority(key, value)`: changes key's value.
- ▶ We'll use a priority queue to hold nodes not yet added to tree.

```
def prim(graph, weight):
    tree = UndirectedGraph()

    estimated_predecessor = {node: None for node in graph.nodes}
    cost = {node: float('inf') for node in graph.nodes}
    priority_queue = PriorityQueue(cost)

    while priority_queue:
        u = priority_queue.extract_min()
        if estimated_predecessor[u] is not None:
            tree.add_edge(estimated_predecessor[u], u)
        for v in graph.neighbors(u):
            if weight(u, v) < cost[v] and v not in tree.nodes:
                priority_queue.decrease_priority(v, weight(u, v))
                cost[v] = weight(u, v)
                estimated_predecessor[v] = u
    return tree
```

# Prim and Dijkstra

- ▶ This is a lot like Dijkstra's Algorithm for s.p.d.!
- ▶ Both: at each step, extract node with smallest cost, update its edges. (Prim: only those edges to nodes not in tree).
- ▶ Dijkstra update of  $(u, v)$ :

$$\text{cost}[v] = \min(\text{cost}[v], \text{cost}[u] + \text{weight}(u, v))$$

- ▶ Prim update of  $(u, v)$ :

$$\text{cost}[v] = \min(\text{cost}[v], \text{weight}(u, v))$$

# DSC 40B

Theoretical Foundations II

## Time Complexity

# Time Complexity

- ▶ A priority queue can be implemented using a **heap**.
- ▶ If a **binary min-heap** is used:
  - ▶ `PriorityQueue(est)` takes  $\Theta(V)$  time.
  - ▶ `.extract_min()` takes  $O(\log V)$  time.
  - ▶ `.decrease_priority()` takes  $O(\log V)$  time.

# Time Complexity

```
def prim(graph, weight):
    tree = UndirectedGraph()

    estimated_predecessor = {node: None for node in graph.nodes}
    cost = {node: float('inf') for node in graph.nodes}
    priority_queue = PriorityQueue(cost)

    while priority_queue:
        u = priority_queue.extract_min()
        if estimated_predecessor[u] is not None:
            tree.add_edge(estimated_predecessor[u], u)
        for v in graph.neighbors(u):
            if weight(u, v) < cost[v] and v not in tree.nodes:
                priority_queue.decrease_priority(v, weight(u, v))
                cost[v] = weight(u, v)
                estimated_predecessor[v] = u

    return tree
```

# Time Complexity

- ▶ Using a **binary heap**...
- ▶ Overall:  $\Theta(V \log V + E \log V)$ .
- ▶ Since graph is assumed connected,  $E = \Omega(V)$ .
- ▶ So this simplifies to  $\Theta(E \log V)$ .

# Fibonacci Heaps

- ▶ A priority queue can be implemented using a **heap**.
- ▶ If a **Fibonacci min-heap** is used:
  - ▶ `PriorityQueue(est)` takes  $\Theta(V)$  time.
  - ▶ `.extract_min()` takes  $\Theta(\log V)$  time<sup>1</sup>.
  - ▶ `.decrease_priority()` takes  $O(1)$  time.

---

<sup>1</sup>Amortized.

# Time Complexity

```
def prim(graph, weight):
    tree = UndirectedGraph()

    estimated_predecessor = {node: None for node in graph.nodes}
    cost = {node: float('inf') for node in graph.nodes}
    priority_queue = PriorityQueue(cost)

    while priority_queue:
        u = priority_queue.extract_min()
        if estimated_predecessor[u] is not None:
            tree.add_edge(estimated_predecessor[u], u)
        for v in graph.neighbors(u):
            if weight(u, v) < cost[v] and v not in tree.nodes:
                priority_queue.decrease_priority(v, weight(u, v))
                cost[v] = weight(u, v)
                estimated_predecessor[v] = u

    return tree
```

# Time Complexity

- ▶ Using a **Fibonacci heap**...
- ▶ Overall:  $\Theta(V \log V + E)$ .

# Fibonacci vs. Binary Heaps

- ▶ Using Fibonacci heaps improves time complexity when graph is dense.
- ▶ E.g., if  $E = \Theta(V^2)$ :
  - ▶ Prim's with Fibonacci:  $\Theta(E) = \Theta(V^2)$
  - ▶ Prim's with binary:  $\Theta(E \log E) = \Theta(V^2 \log V)$ .
- ▶ But Fibonacci heaps are **hard** to implement; have large constants.
- ▶ Binary heaps used more in practice despite complexity.

# DSC 40B


Theoretical Foundations II

## Correctness of Prim's Algorithm

# Being Greedy

- ▶ At every step, we add the lightest edge.
- ▶ Is this “safe”?
- ▶ Yes! This is guaranteed to find an MST.

# Promising Subtrees




- ▶ Let  $G = (V, E, \omega)$  be a weighted graph.
- ▶ A subgraph  $T' = (V', E')$  is **promising** if it is “part” of some MST.
  - ▶ That is, it is an “MST in progress”
  - ▶ Not necessarily a tree!
- ▶ That is, there exists an MST  $T = (V, E_{\text{mst}})$  such that  $E' \subset E_{\text{mst}}$ .
- ▶ Hint: a “promising subtree” where  $V' = V$  is an MST!

## Main Idea

Prim's starts with a promising subtree  $T$ . At each step, adds lightest edge from a node within  $T$  to a node outside of  $T$ .


We'll show each new edge results in a larger promising subtree. Eventually the promising subtree becomes a full MST.

# Claim



- ▶ Let  $G = (V, E, \omega)$  be a weighted graph.
- ▶ Suppose  $T' = (V', E')$  is a promising subtree for an MST of  $G$ .
- ▶ Let  $e = (u, v)$  be a **lightest edge** from a node in  $T'$  to a node outside of  $T'$ . (Prim).
- ▶ Then adding  $(u, v)$  to  $T'$  results in another **promising subtree**.

# Proof



- ▶ Suppose  $T_{\text{mst}}$  is an MST that includes  $T'$ .
- ▶ If  $T_{\text{mst}}$  includes  $e$ , we're done:  $T' + e$  is promising.
- ▶ If it doesn't include  $e$ , it must have an edge  $f$  that connects  $T'$  to rest of the graph.
- ▶ Swap  $f$  with  $e$  in  $T_{\text{mst}}$ . The result is a tree, and it must be a MST since  $\omega(e) \leq \omega(f)$ .
- ▶ So there is an MST that contains  $T' + e$ .