

DSC40B: Theoretical Foundations of Data Science II

Lecture 16: *Minimum Spanning Tree,
properties, and general greedy algorithms*

Instructor: Yusu Wang

Previously

- ▶ Given directed or undirected graphs
 - ▶ Graph search / traversal strategies (DFS / BFS)
 - ▶ Single source shortest paths in weighted graphs
 - ▶ Bellman-Ford algorithm for general graphs
 - ▶ Dijkstra algorithm for graphs with positive edge weights
- ▶ Today:
 - ▶ Computing a minimum spanning tree (**MST**) of **an undirected graph**



Trees, spanning trees, and minimum spanning tree



Trees

- ▶ An undirected graph $G = (V, E)$ is a **tree** if and only if
 - ▶ (i) it is connected; and
 - ▶ (ii) it is acyclic (i.e., does not contain any cycle)

- ▶ Claim [Tree Edges]:
 - ▶ If $T = (V, E)$ is a tree, then we have that $|E| = |V| - 1$

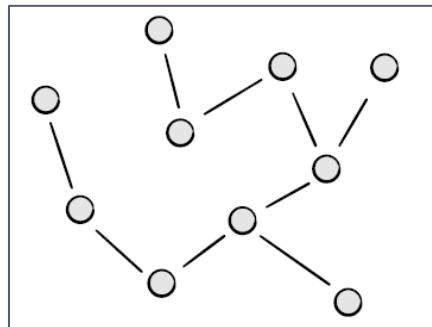


Alternative definition

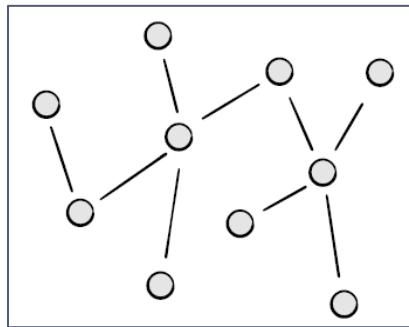
▶ Alternative definition:

- ▶ An undirected graph $G = (V, E)$ is a **tree** if and only if that (i) it is connected; and (ii) $|E| = |V| - 1$

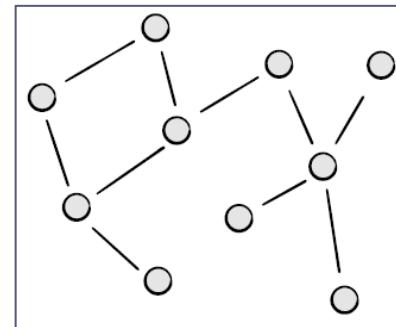
A tree



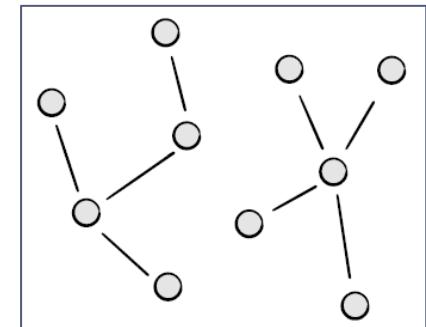
A tree



NOT a tree

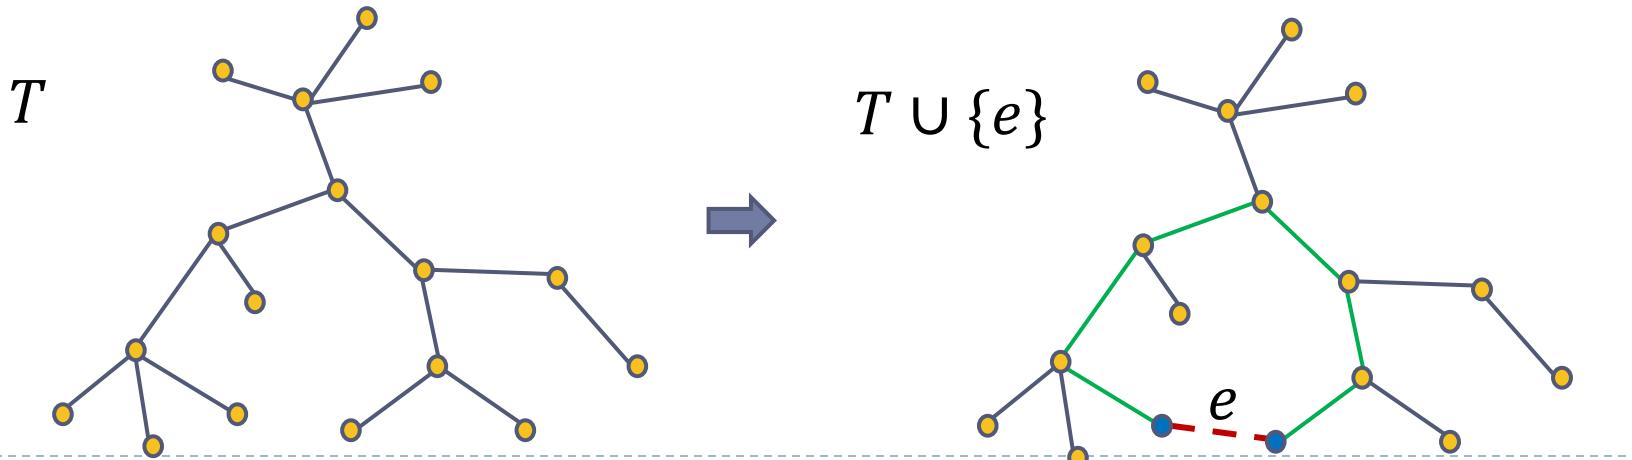


NOT a tree



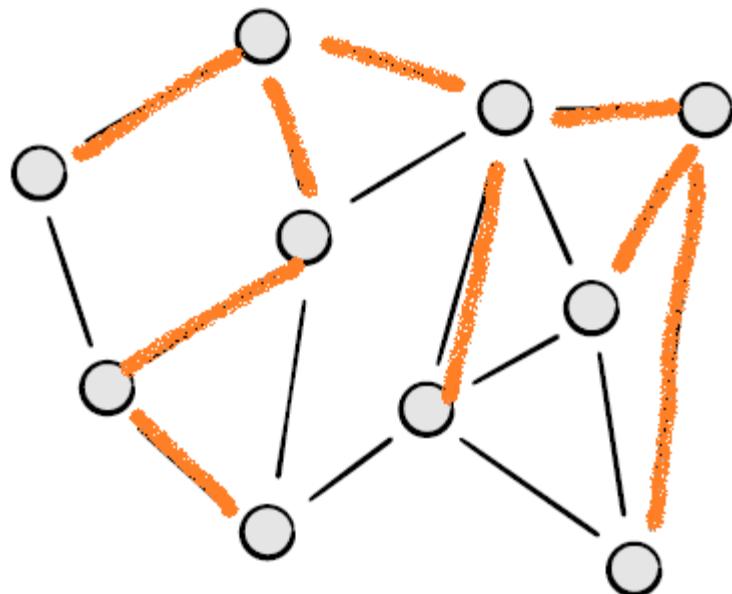
Remarks

- ▶ Key properties: If $T = (V, E)$ is a tree,
 - ▶ there is a unique path between any two nodes in V
 - ▶ adding any other edge e to T will create a unique cycle containing e
 - ▶ i.e., $T \cup \{e\}$ contains a cycle for any $e \notin T$
 - ▶ removing an edge from T will disconnect it
- ▶ Out of all connected graphs on n nodes, a tree has least number (i.e., $n - 1$) of edges



Spanning Tree

- Given an undirected graph $G = (V, E)$, a **spanning tree of G** is any graph $T = (V, E' \subseteq E)$ that is a tree.



Example of spanning trees for the graph on the right.

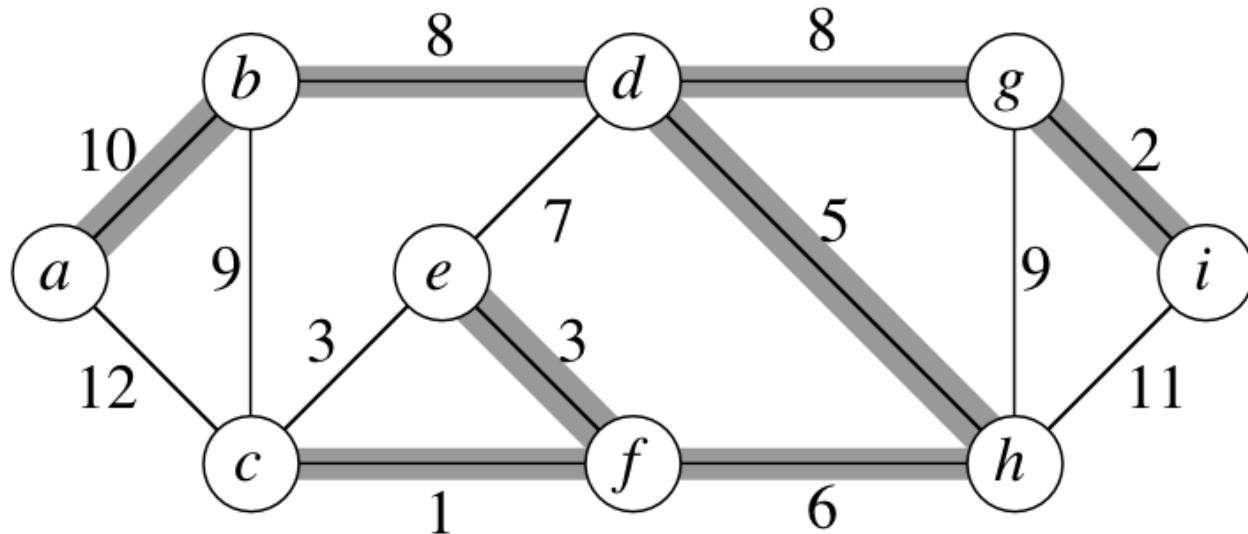
Spanning Tree

- ▶ Given an undirected graph $G = (V, E)$, a **spanning tree of G** is any graph $T = (V, E' \subseteq E)$ that is a tree.
- ▶ Intuitively, a spanning tree of G contains smallest number of edges in E to connect all nodes in G .
- ▶ Note that if the input graph G is not connected, then there exists no spanning tree.
 - ▶ We can talk about spanning forest, consisting a set of spanning trees, one for each connected component in G .



Minimum spanning tree (MST)

- ▶ Weight of spanning tree T of a weighted graph $G = (V, E)$ is
 - ▶ the total weights of all edges in T , i.e., $\omega(T) = \sum_{e \in T} \omega(e)$,
 - ▶ where $\omega: E \rightarrow \mathbb{R}$ is the edge weights associated to G .
- ▶ A minimum spanning tree (MST) of a weighted graph $G = (V, E)$ is a spanning tree with smallest possible weight.



Weight of this spanning tree: 43

Turns out this is also a minimum spanning tree.

MSTs

- ▶ MST may not be unique
- ▶ All MSTs of a given graph $G = (V, E)$ have the same number of edges!
 - ▶ They all have $|V| - 1$ number of edges
- ▶ If all edges in input graph have the same weight, then how can we find a MST for it?
 - ▶ Any spanning tree of it is a minimum spanning tree!

Exercise:

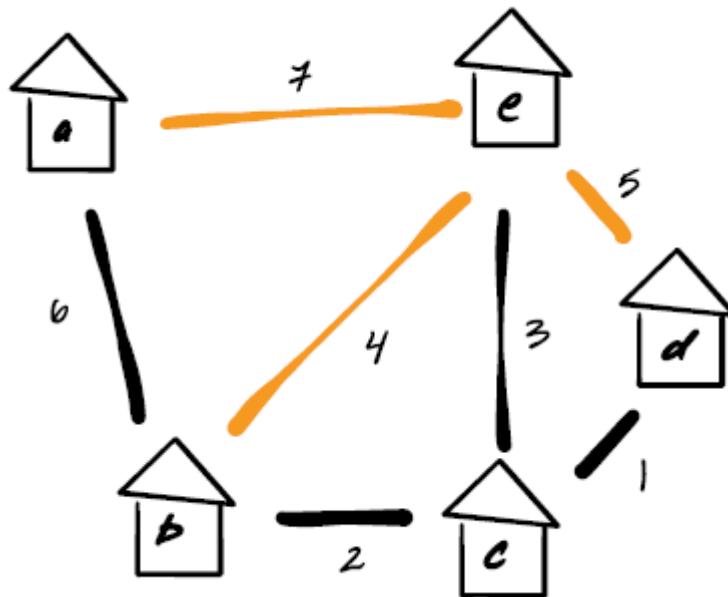
Design an algorithm to compute an MST for a graph where all edges have weight 1



Motivation and properties of MST



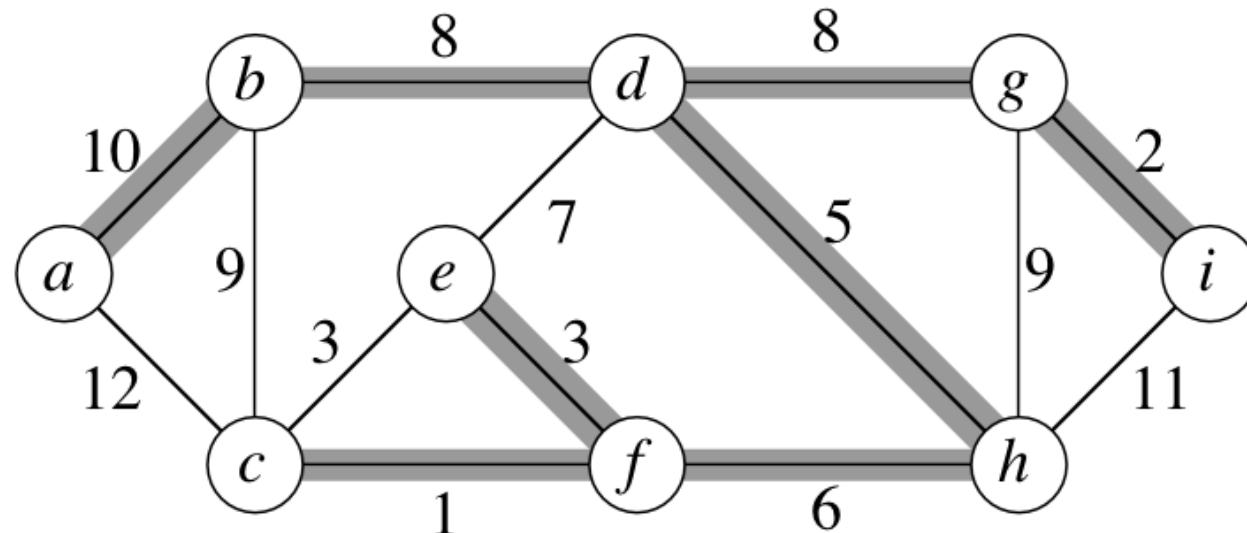
Motivation



- ▶ Among all possible road segment choices, build a set of road segments so that all houses are connected and the total cost is minimized
- ▶ **Solution:** Find the MST of the input weighted graph where edge weight represents the cost of build that road segment.

MST Problem

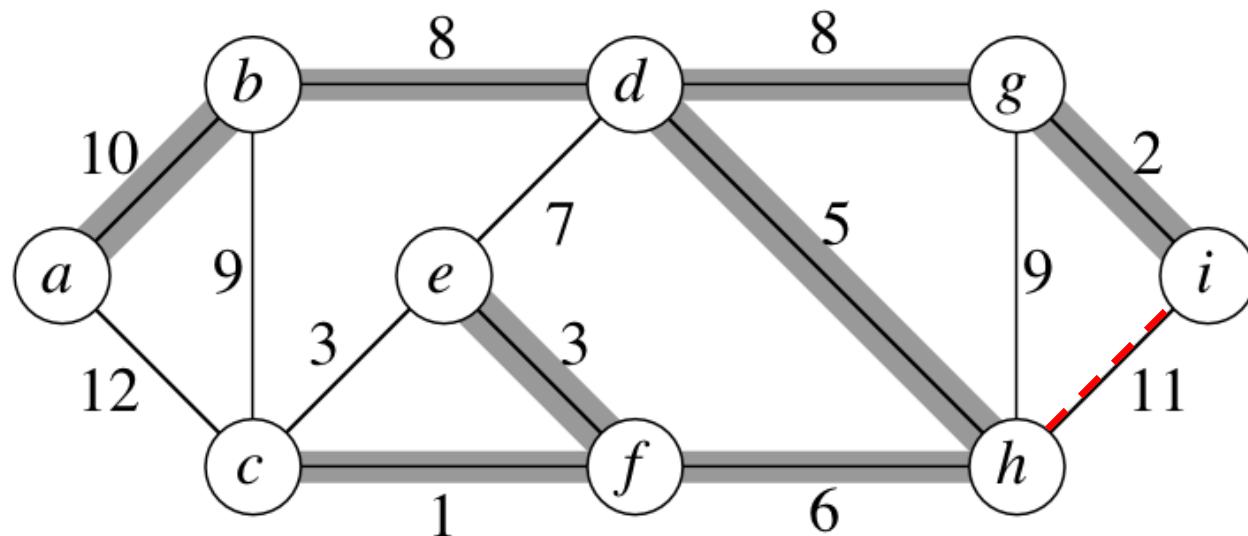
- ▶ **Input:**
 - ▶ a weighted undirected graph G
- ▶ **Output:**
 - ▶ the set of edges in a MST of G



Key property

Key property of MST:

- Given a MST T of $G = (V, E)$, let $e \in E$ be any edge in E but not in T . The following then holds:
 - there is a unique cycle C containing e in $T \cup \{e\}$.
 - e has the largest weight among all edges in this cycle C .



Key property

Key property of MST:

- ▶ Given a MST T of $G = (V, E)$, let $e \in E$ be any edge in E but not in T . The following then holds:
 - ▶ there is a unique cycle C containing e in $T \cup \{e\}$.
 - ▶ e is an edge with largest weight in C .

- ▶ Proof sketch:
 - ▶ If e does not have largest weight, let $e' \in C$ be an edge with largest weight in C .
 - ▶ $T' = T - \{e'\} + \{e\}$ is also a spanning tree of G
 - ▶ $\text{weight}(T') \leq \text{weight}(T) \Rightarrow T$ cannot be MST.
 - ▶ Contradiction $\Rightarrow e$ must have largest weight in C .



First greedy algorithm for MST: Prim's algorithm

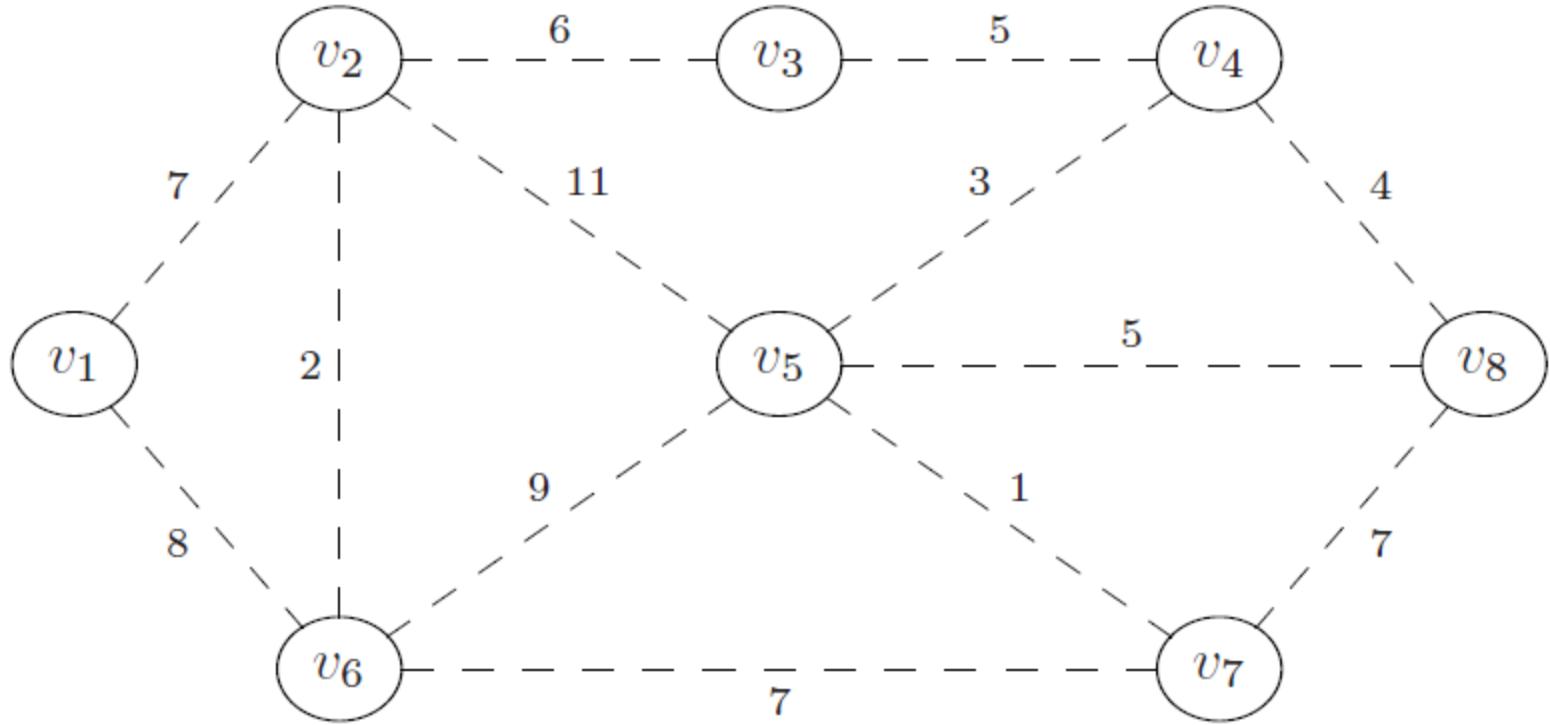


General greedy idea:

- ▶ **Input:**
 - ▶ a weighted undirected graph $G = (V, E)$, with $\omega: E \rightarrow R$
- ▶ **Output:**
 - ▶ the set of edges in a MST T of G
- ▶ A MST T is $V - 1$ number of edges that connect all nodes, with no cycle.
- ▶ Intuitively, we will grow the tree edge-by-edge, and choose “**safe**” edges greedily to incrementally build T
 - ▶ such that any time, the edges we choose will form a part of some MST



Example



- ▶ What is a “**safe**” edge to add first?
-



Two greedy algorithms

- ▶ Two greedy algorithms
 - ▶ Today: **Prim's** algorithm
 - ▶ Next class: **Kruskal's** algorithm
 - ▶ They differ in the order of edges they visit and thus ``safe'' edges they add



Idea for Prim's algorithm

- ▶ **Input:**
 - ▶ a weighted undirected graph $G = (V, E)$, with $\omega: E \rightarrow R$
- ▶ **Output:**
 - ▶ the set of edges in a MST T of G

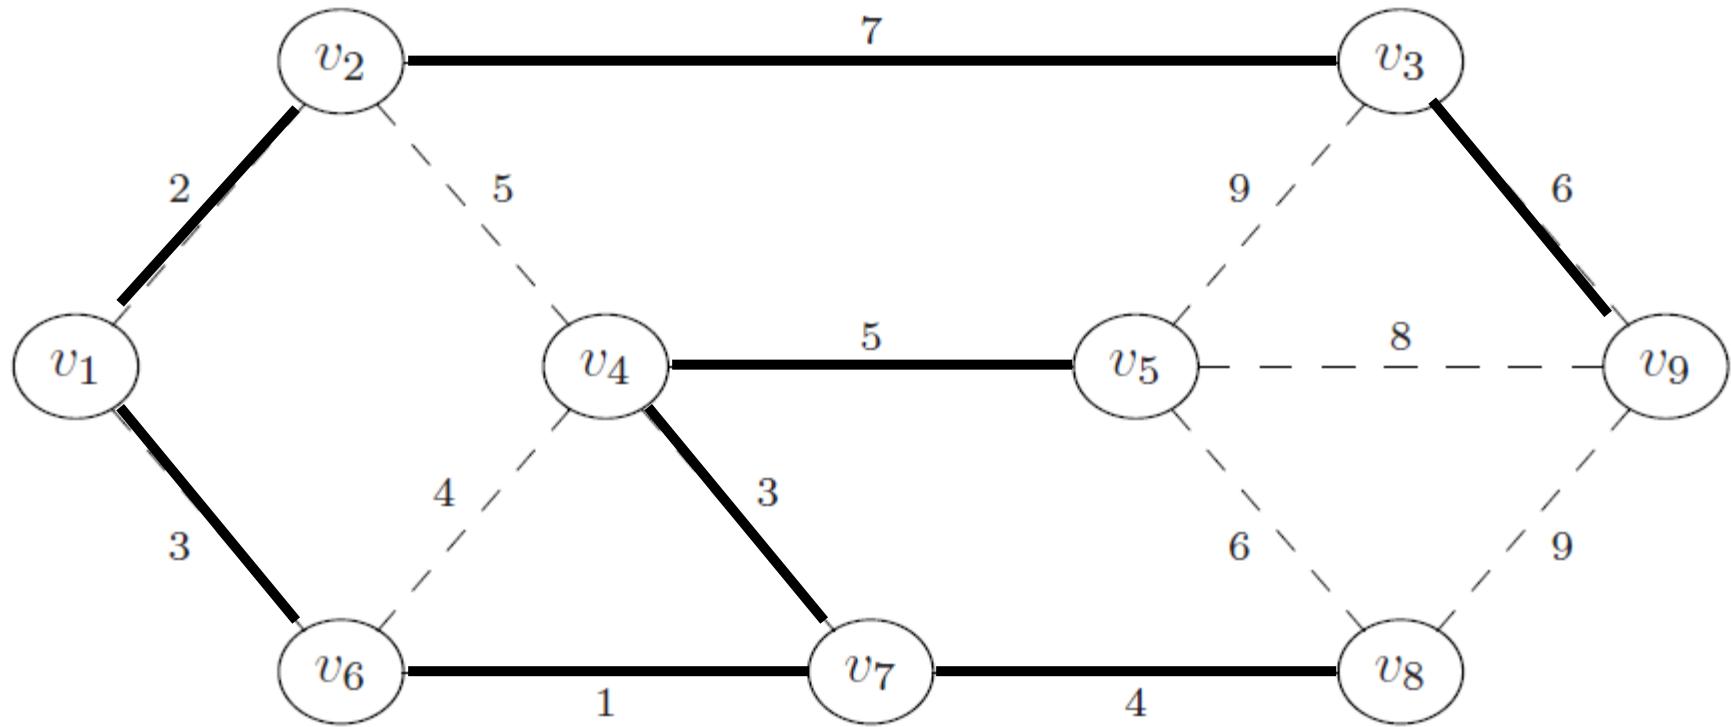
- ▶ **Intuitively,**
 - ▶ Incrementally grow a partial tree $T(S) \subseteq E$ connecting a subset of nodes $S \subset V$
 - ▶ At the beginning of each iteration, $T(S)$ is a sub-tree of **some** MST of G
 - ▶ At each iteration, grow $T(S')$ to include one more vertex $S' = S \cup \{u\}$
 - ▶ such that $T(S')$ is still a sub-tree of **some** MST of G
 - ▶ the new node is reached via a greedy choice of a **crossing-edge**
 - ▶ in particular, the greedy choice is the minimum weight edge connect some node in S to some node in $V = S$ (i.e., outside S)

High level outline (not code)

```
procedure PrimMST(G)
1  $U \leftarrow V(G) - \{v_1\}$ ; /*  $V(G)$  = set of vertices of graph  $G$  */
2  $v_1.parent \leftarrow \text{NULL}$ ;
3 while ( $U \neq \emptyset$ ) and ( $\exists$  edge from  $(V(G) - U)$  to  $U$ ) do
4    $(v_i, v_j) \leftarrow$  minimum weight edge from  $V(G) - U$  to  $U$ ;
5    $v_j.parent \leftarrow v_i$ ;
6    $U \leftarrow U - \{v_j\}$ ;
7 end
```

- ▶ U : unconnected vertices
- ▶ $S = V - U$: vertices connected by current partial tree

Example



- ▶ Suppose we grow the tree starting from v_1

Correctness

MST Theorem:

Let T be a sub-tree of a minimum spanning tree.

If e is a minimum weight edge connecting T to some vertex not in T , then $T \cup \{e\}$ is a subtree of a minimum spanning tree.

- ▶ Key to the correctness of PrimMST algorithm.
 - ▶ Invariance:

each time PrimMST() algorithm grows the partial tree (i.e, adds another edge to it), the invariance is that the new tree is still a subtree of **some** minimum spanning tree of input graph G .
 - ▶ Termination:

when all nodes are connected, we obtain a MST of G .
(or if we cannot reach all nodes, then the input graph is not connected)

Implementation of Prim's algorithm



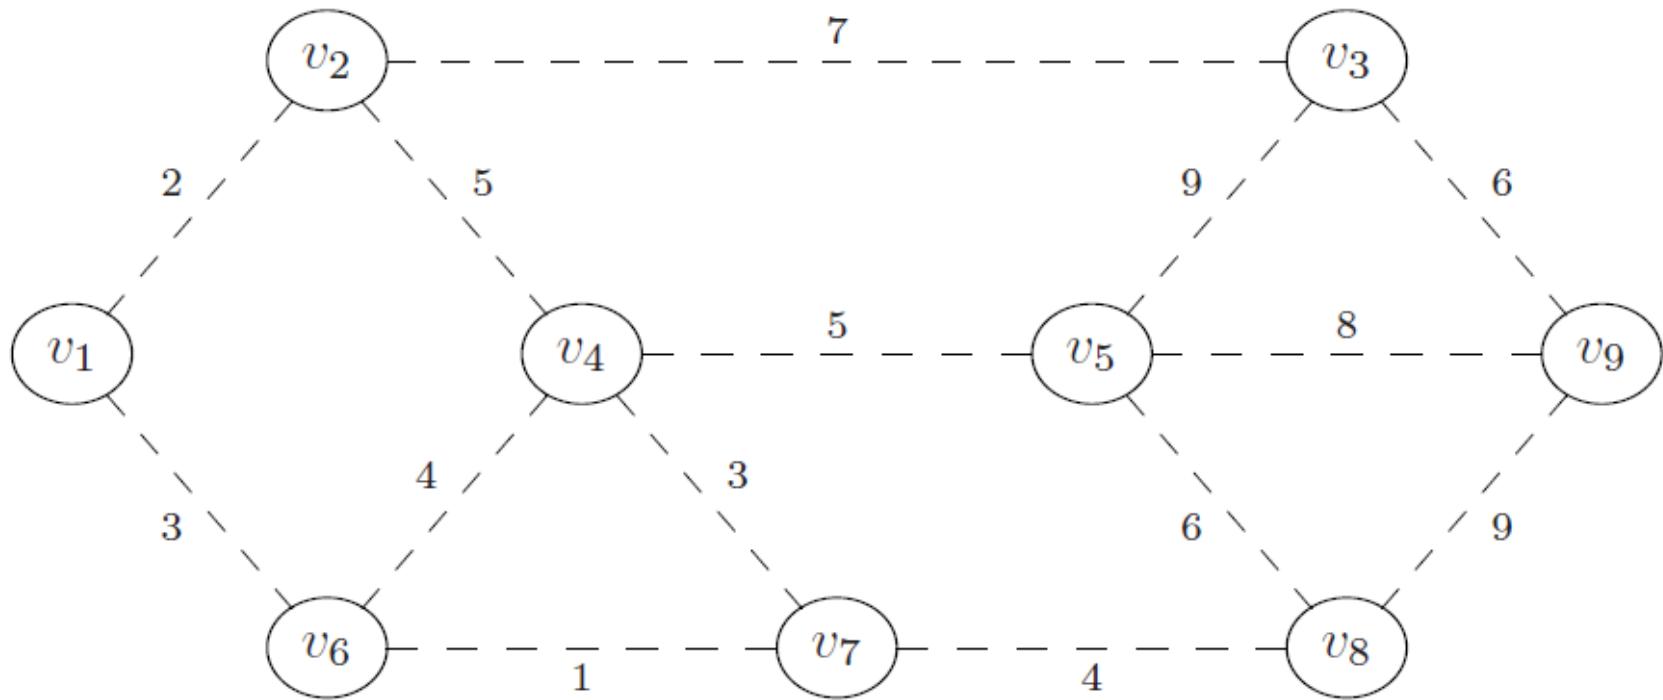
Naïve implementation of Prim's Alg

```
procedure PrimMST(G)
1  $U \leftarrow V(G) - \{v_1\}$ ; /*  $V(G)$  = set of vertices of graph  $G$  */
2  $v_1.\text{predecessor} \leftarrow \text{NULL}$ ;
3 while ( $U \neq \emptyset$ ) and ( $\exists$  edge from  $(V(G) - U)$  to  $U$ ) do
4    $(v_i, v_j) \leftarrow$  minimum weight edge from  $V(G) - U$  to  $U$ ;
5    $v_j.\text{predecessor} \leftarrow v_i$ ;
6    $U \leftarrow U - \{v_j\}$ ;
7 end
```

- ▶ Naïve implementation: linear scan all edges to identify min-weight edge (v_i, v_j) at each iteration
- ▶ Total time complexity: $O(VE)$

First improvement

- ▶ Storing costs at nodes
 - ▶ Each unvisited nodes v in U maintain $v.\text{cost}$, which is the smallest weight of any edge from v to visited nodes in S



Outline of first improvement

```
procedure PrimMST(G)
```

- 1 $U \leftarrow V(G)$; */* $V(G)$ = set of vertices of graph G */*
- 2 **foreach** $v_i \in V(G) - \{v_1\}$ **do** $v_i.\text{cost} \leftarrow \infty$;
- 3 $v_1.\text{cost} \leftarrow 0$;
- 4 $v_1.\text{predecessor} \leftarrow \text{NULL}$;
- 5 **while** ($U \neq \emptyset$) **and** ($v_i.\text{cost} < \infty$ for some $v_i \in U$) **do**
- 6 $v_j \leftarrow v_i \in U$ with minimum $v_i.\text{cost}$;
- 7 $U \leftarrow U - \{v_j\}$; */* Remove v_j from U */*
 / $(v_j, v_j.\text{predecessor})$ is an MST edge */*
- 8 **foreach** edge (v_j, v_k) incident on v_j **do**
- 9 **if** (v_k is in U **and** $\text{weight}(v_j, v_k) < v_k.\text{cost}$) **then**
- 10 $v_k.\text{predecessor} \leftarrow v_j$;
- 11 $v_k.\text{cost} \leftarrow \text{weight}(v_j, v_k)$;
- 12 **end**
- 13 **end**
- 14 **end**

Better implementation

- ▶ Similar to Dijkstra algorithm, we can use priority-queue to significantly speed up the time complexity!
- ▶ In particular, we need a data structure to maintain the costs of unvisited nodes, which supports:
 - ▶ deleting the node with minimum cost (`.extract_min !`)
 - ▶ update (decrease) the cost value stored at a node (`change_priority !`)
- ▶ In our case, a priority queue stores `(key, value)` pairs, where key refers to identity of some node, while value is the cost of this node.



Recall Heap implementation

- ▶ A priority queue can be implemented using a (min) heap
- ▶ min-heap implementation of priority queue:
 - ▶ `PriorityQueue(priorities)`: takes $\Theta(n)$ time for $n = |\text{priorities}|$
 - ▶ `.extract_min()` : takes $\Theta(\log n)$ time where n is the size of priority queue
 - ▶ `.change_priority(key, value)` : takes $\Theta(\log n)$ time where n is the size of priority queue



Final implementation of Prim's Alg

```
def prim(graph, weight):
    tree = UndirectedGraph()

    estimated_predecessor = {node: None for node in graph.nodes}
    initial_costs = {node: float('inf') for node in graph.nodes}
    priority_queue = PriorityQueue(costs)

    while priority_queue:
        u = priority_queue.extract_min()
        tree.add_edge(estimated_predecessor[u], u)
        for v in graph.neighbors(u):
            if weight(u, v) < cost[v] and v not in tree.nodes:
                priority_queue.decrease_priority(v, weight(u, v))
                estimated_predecessor[v] = u

    return tree
```



Time complexity analysis

- ▶ We use min-heap to implement the priority queue
- ▶ The maximum size of Q is V
- ▶ # iterations of While-loop?
 - ▶ V
- ▶ # iterations of each call of the inner for-loop?
 - ▶ $\deg(v_j)$
- ▶ Total #times lines 7—10 are executed:
 - ▶ $\sum_{v_j \in V} \deg(v_j) = 2E$

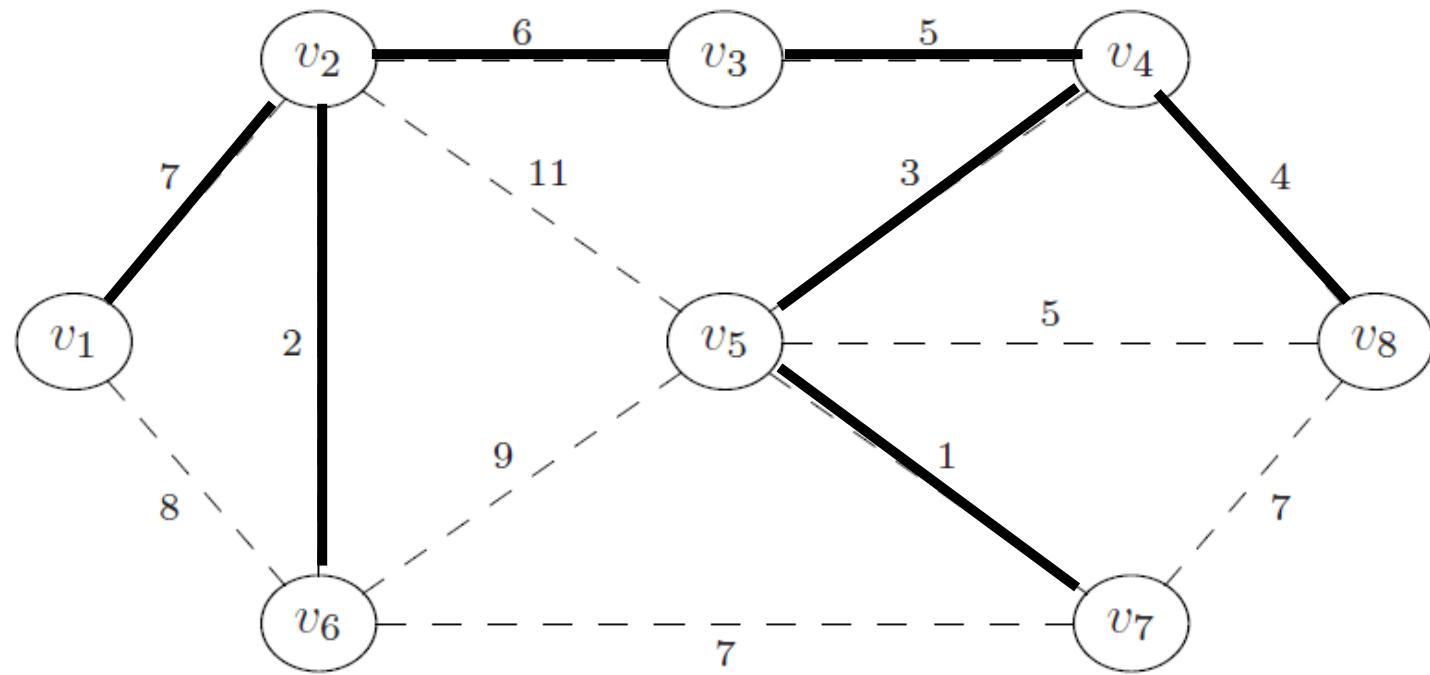


- ▶ **Initialize priority_queue**
 - ▶ Total cost: $\Theta(V)$
- ▶ **decrease_priority**
 - ▶ Total #: at most $2E$
 - ▶ Total cost: $O(E \lg V)$
- ▶ **extract_min**
 - ▶ Total #: V
 - ▶ Total cost: $\Theta(V \lg V)$

Total time complexity:
 $\Theta((V + E) \lg V)$



Example



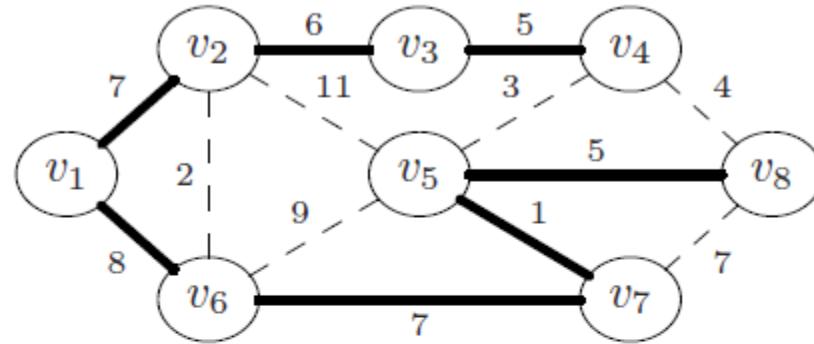
Comparison with Dijkstra algorithm

- ▶ **Dijkstra:**
 - ▶ Each node maintains the best distance estimate from source to the current node
- ▶ **Prim's:**
 - ▶ Each node (not yet visited) maintains the minimum weight of any edge to reach a visited-node.

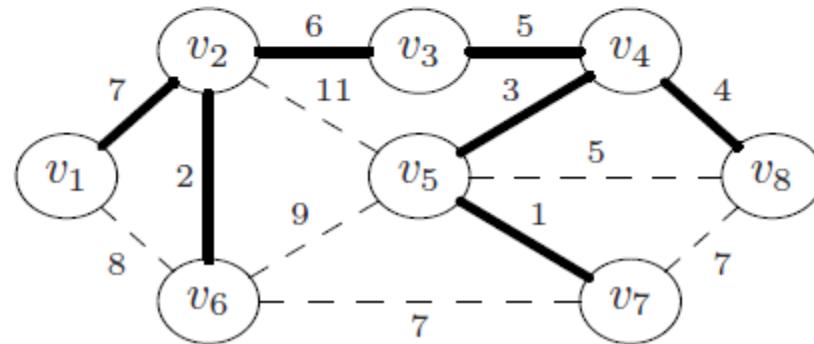


Comparison with Dijkstra

Shortest Path Tree:



Minimum Spanning Tree:



Summary and remarks

- ▶ **Prim's algorithm:**
 - ▶ A greedy algorithm which repeatedly choose the minimum-weight edge to reach an unvisited node
 - ▶ Share similarity to Dijkstra algorithm
 - ▶ Runs in $\Theta((V + E) \lg V)$ time using min-heap
- ▶ Similar to Dijkstra algorithm, we can further improve the time complexity to $\Theta(E + V \lg V)$ using Fibonacci heap, which is a more efficient implementation of priority queue.
- ▶ **Next time,**
 - ▶ Another greedy algorithm, called Kruskal algorithm, which has other properties too.



FIN

